



Is this your card?
The Magic of Computer Vision

IANADS

- @guyroyse (probably)



@cr0wst

IANAPD



@cr0wst



@cr0wst

"Getting someone to pick a card is almost impossible. Go to a party, fan out a deck of cards, and ask someone to pick a card.

Everyone will say no."

-Penn Jillette



@cr0wst



@cr0wst

Every magic trick consists
of three parts, or acts.

-The Prestige



@cr0wst

The first part is called the pledge, the magician shows you something ordinary.

-The Prestige



@cr0wst

The second act is called the turn, the magician takes the ordinary something and makes it into something extraordinary.

-The Prestige



@cr0wst

...making something disappear isn't enough; you have to bring it back. That's why every magic trick has a third act, the hardest part, the part we call "The Prestige".

-The Prestige



@cr0wst

Where is the real magic?



@cr0wst

Computer Vision aims to train
computers to understand the visual
world.



@cr0wst

We're going to aim to answer
two questions.



@cr0wst

Where is the card in the video?



@cr0wst

Which card is it?



@cr0wst

```
import cv2
import imutils

def main():
    cap = cv2.VideoCapture(1)
    while True:
        ret, frame = cap.read()
        frame = imutils.resize(frame, 640)
        cv2.imshow("Image", frame)

        if cv2.waitKey(1) & 0xFF == ord('q'):
            break

    cap.release()
    cv2.destroyAllWindows()

main()
```



@cr0wst

```
import cv2
import imutils

def main():
    cap = cv2.VideoCapture(0)
    while True:
        ret, frame = cap.read()
        frame = imutils.resize(frame, 640)
        cv2.imshow("Image", frame)
        if cv2.waitKey(1) & 0xFF == ord('q'):
            break
    cv2.destroyAllWindows()

main()
```

Open Computer Vision Library

import cv2 ←

import imutils →

Convenience functions for basic image processing.



@cr0wst

```
import cv2
import imutils

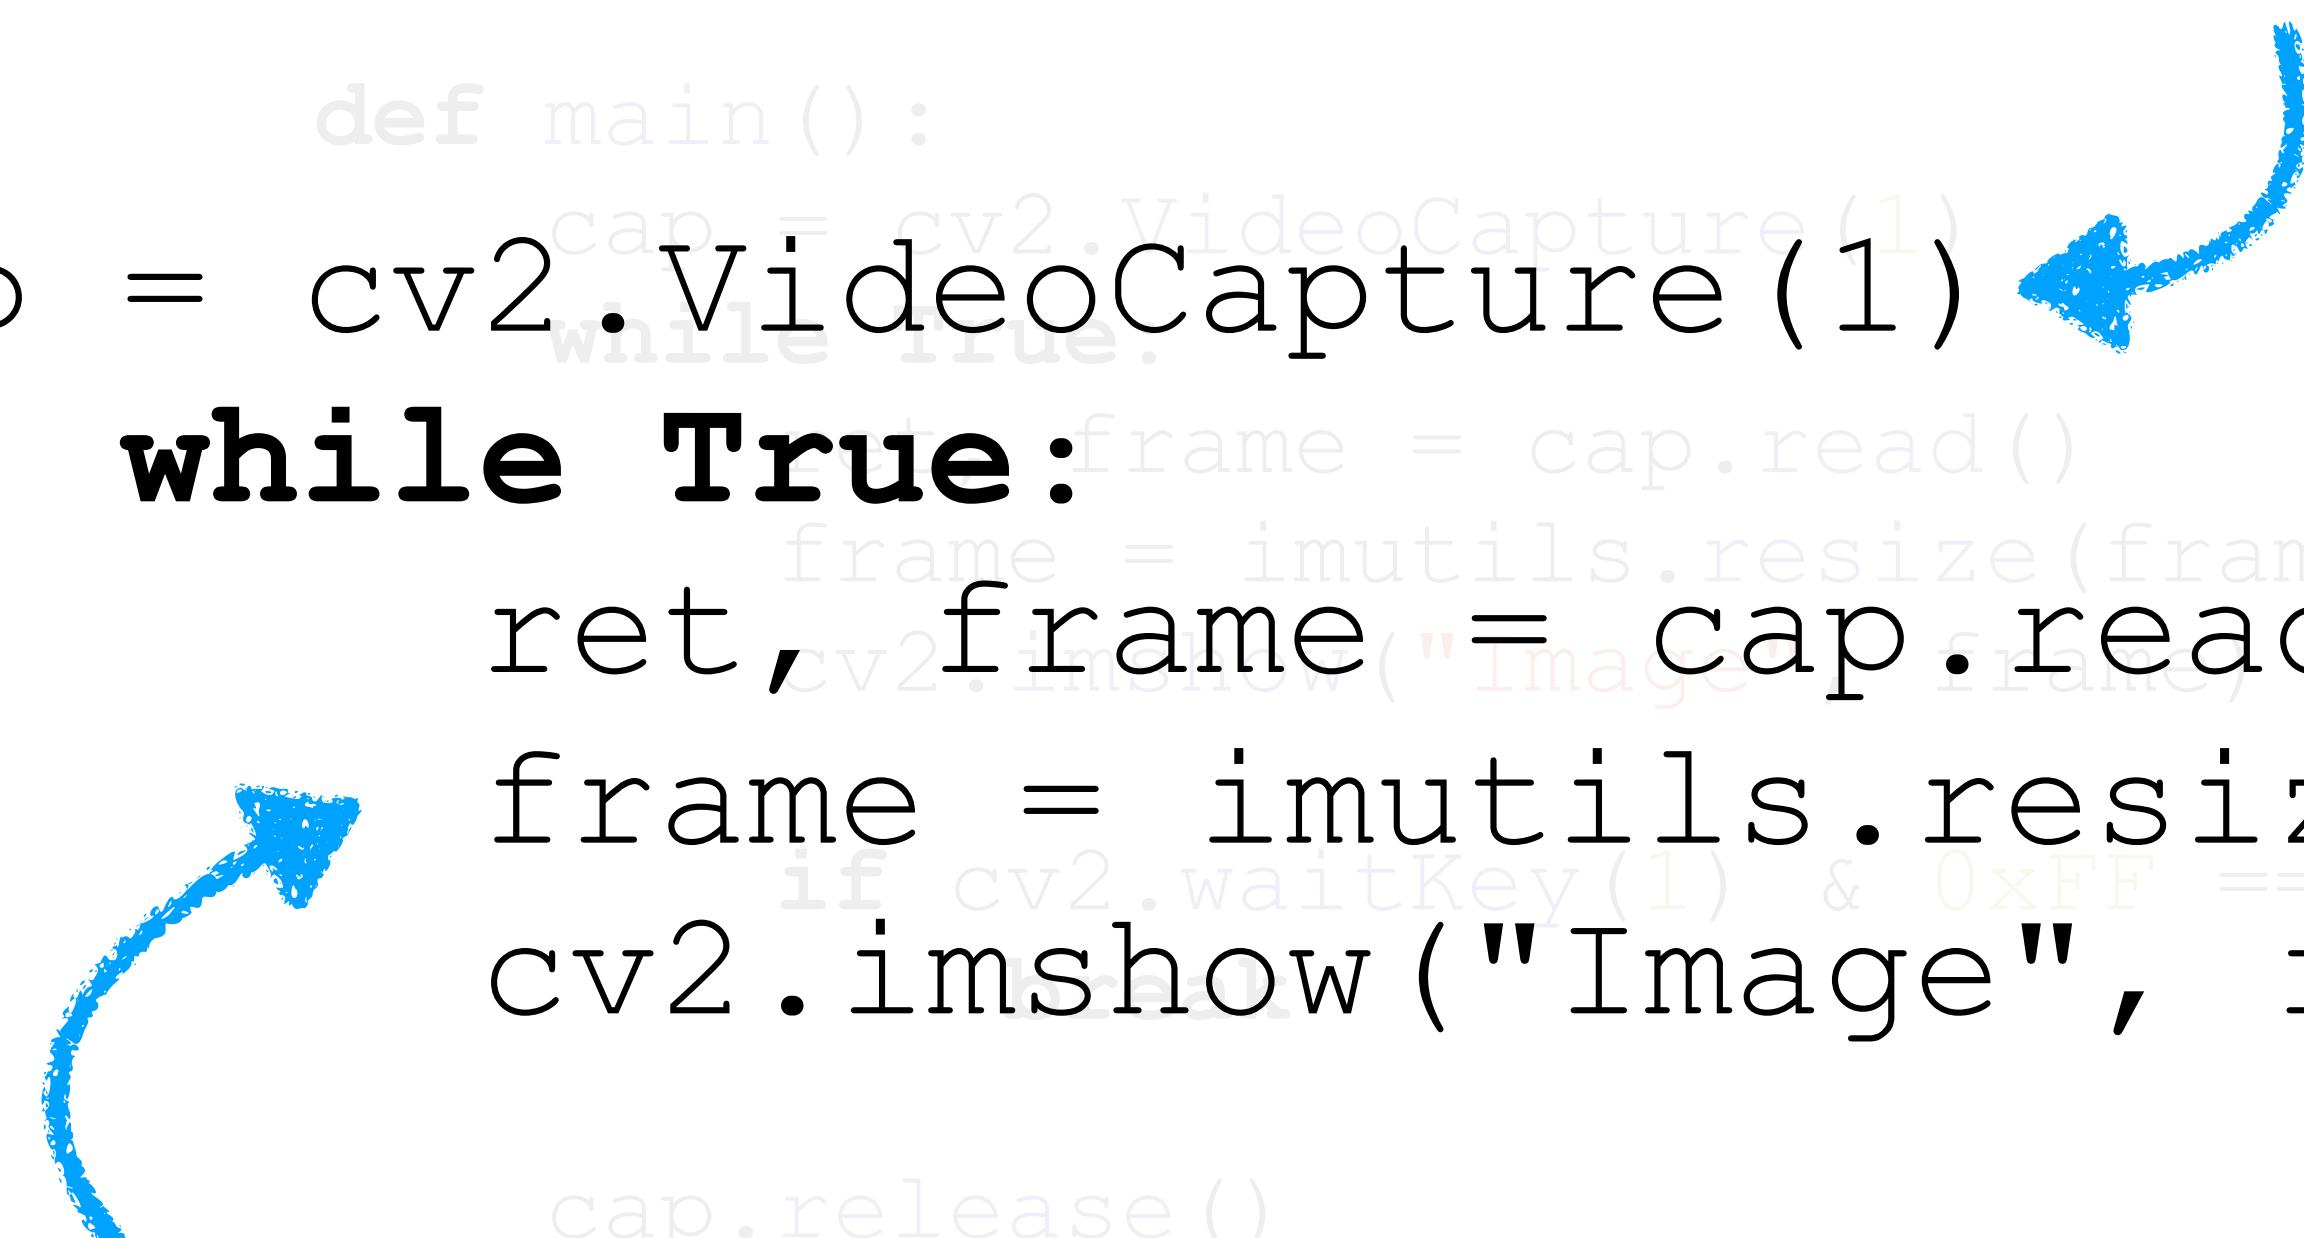
def main():
    cap = cv2.VideoCapture(1)
    while True:
        ret, frame = cap.read()
        frame = imutils.resize(frame, 640)
        cv2.imshow("Image", frame)

    cap.release()
    cv2.destroyAllWindows()

main()
```

Capture the video

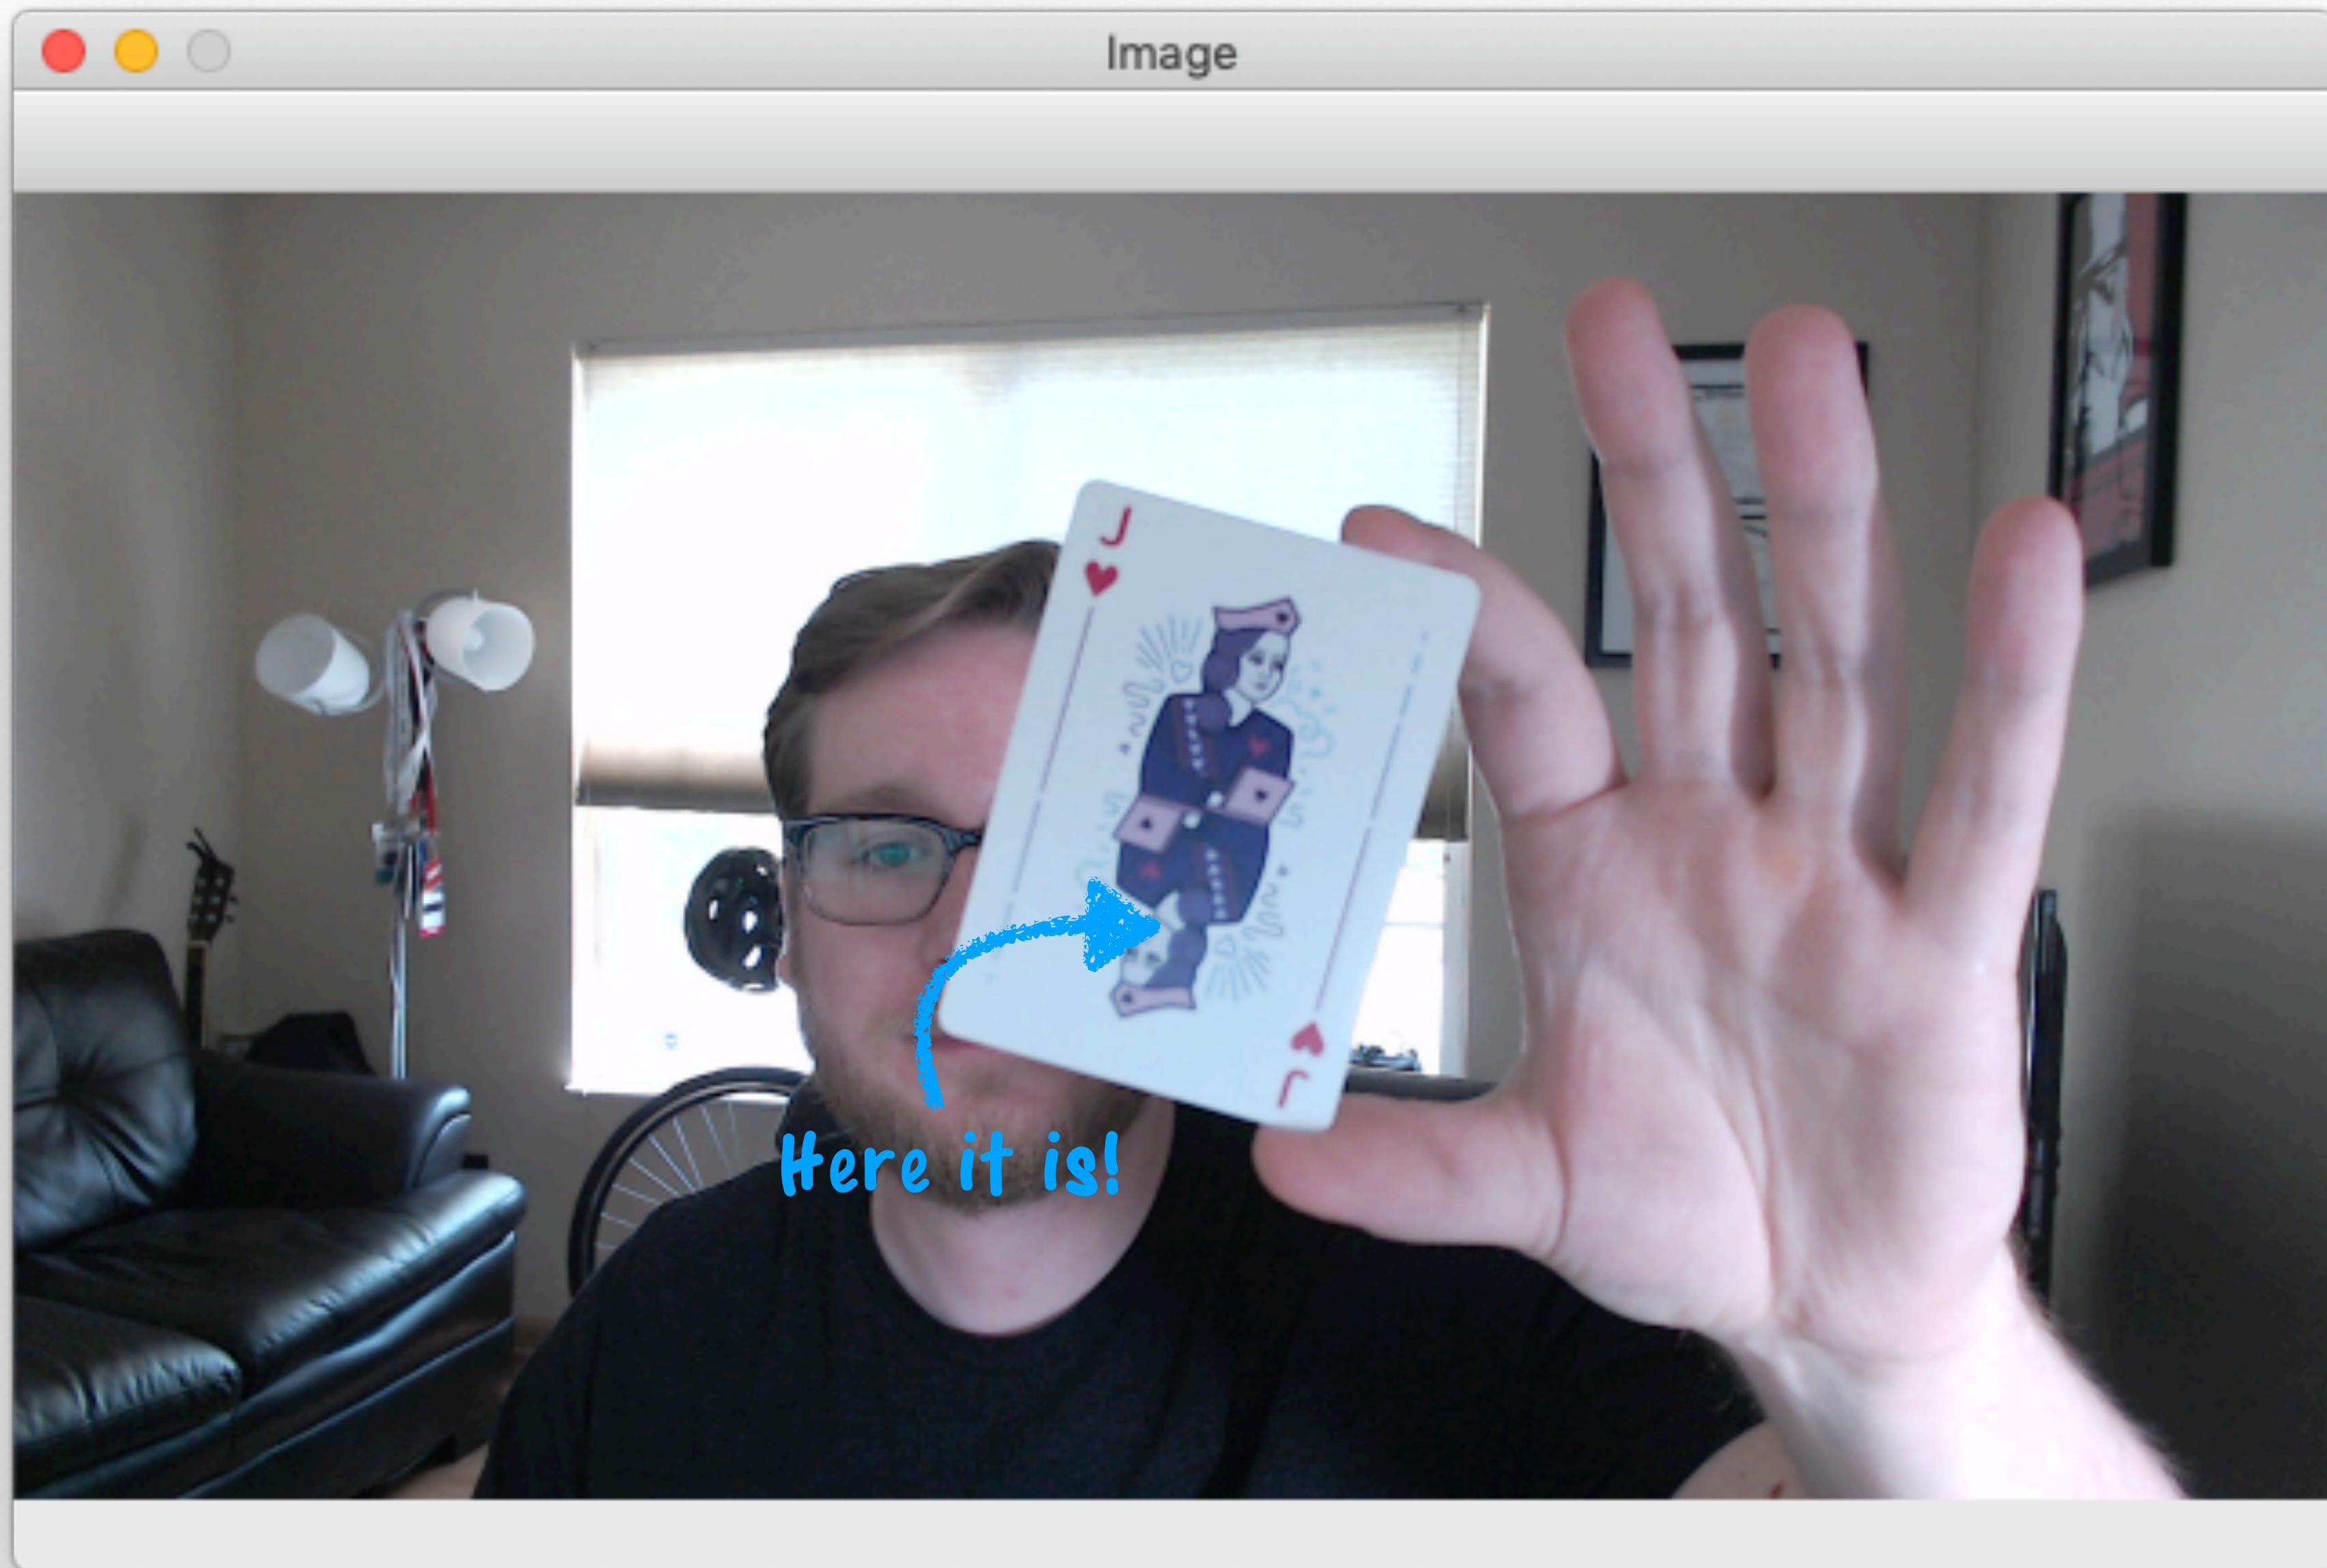
Display each frame as a continuous video loop.



@cr0wst



@cr0wst



@cr0wst

Computers need help.



@cr0wst



@cr0wst

```
import numpy as np

def normalize(frame):
    frame = imutils.resize(frame, 640)
    gray = cv2.cvtColor(frame, cv2.COLOR_BGR2GRAY)
    blur = cv2.bilateralFilter(gray, 10, 15, 15)

    edges = cv2.Canny(blur, 50, 150, True)
    kernel = np.ones((5, 5), np.uint8)
    dilate = cv2.dilate(edges, kernel, iterations=1)

    return dilate
```



@cr0wst

```
import numpy as np

def normalize(frame):
    frame = imutils.resize(frame, 640)
    gray = cv2.cvtColor(frame, cv2.COLOR_BGR2GRAY)
    blur = cv2.bilateralFilter(gray, 10, 15, 15)
import numpy as np
    edges = cv2.Canny(blur, 50, 150, True)
    kernel = np.ones((5, 5), np.uint8)
    dilate = cv2.dilate(edges, kernel, iterations=1)
```

Adds arrays, matrices, large number support, and math functions.



@cr0wst

```
import numpy as np

def normalize(frame):
    frame = imutils.resize(frame, 640)
    gray = cv2.cvtColor(frame, cv2.COLOR_BGR2GRAY)
    blur = cv2.bilateralFilter(gray, 10, 15, 15)
    edges = cv2.Canny(blur, 50, 150, True)
    kernel = np.ones((5, 5), np.uint8)
    dilate = cv2.dilate(edges, kernel, iterations=1)

return dilate
```

This flag is the type of conversion



@cr0wst



@cr0wst

```
import numpy as np  
  
def normalize(frame):  
    frame = imutils.resize(frame, 640)  
    gray = cv2.cvtColor(frame, cv2.COLOR_BGR2GRAY)  
    blur = cv2.bilateralFilter(gray, 10, 15, 15)  
    edges = cv2.Canny(blur, 50, 150, True)  
    kernel = np.ones((5, 5), np.uint8)  
    dilate = cv2.dilate(edges, kernel, iterations=1)  
  
    return dilate
```

This is the sigma space

This is the size of the filter



@cr0wst



@cr0wst

```
import numpy as np

def normalize(frame):
    frame = imutils.resize(frame, 640)
    gray = cv2.cvtColor(frame, cv2.COLOR_BGR2GRAY)
    blur = cv2.bilateralFilter(gray, 10, 17, 15)
edges = cv2.Canny(blur, 50, 150, True)
edges = cv2.Canny(blur, 50, 15, True)
kernel = np.ones((5, 5), np.uint8)
dilate = cv2.dilate(edges, kernel, iterations=1)

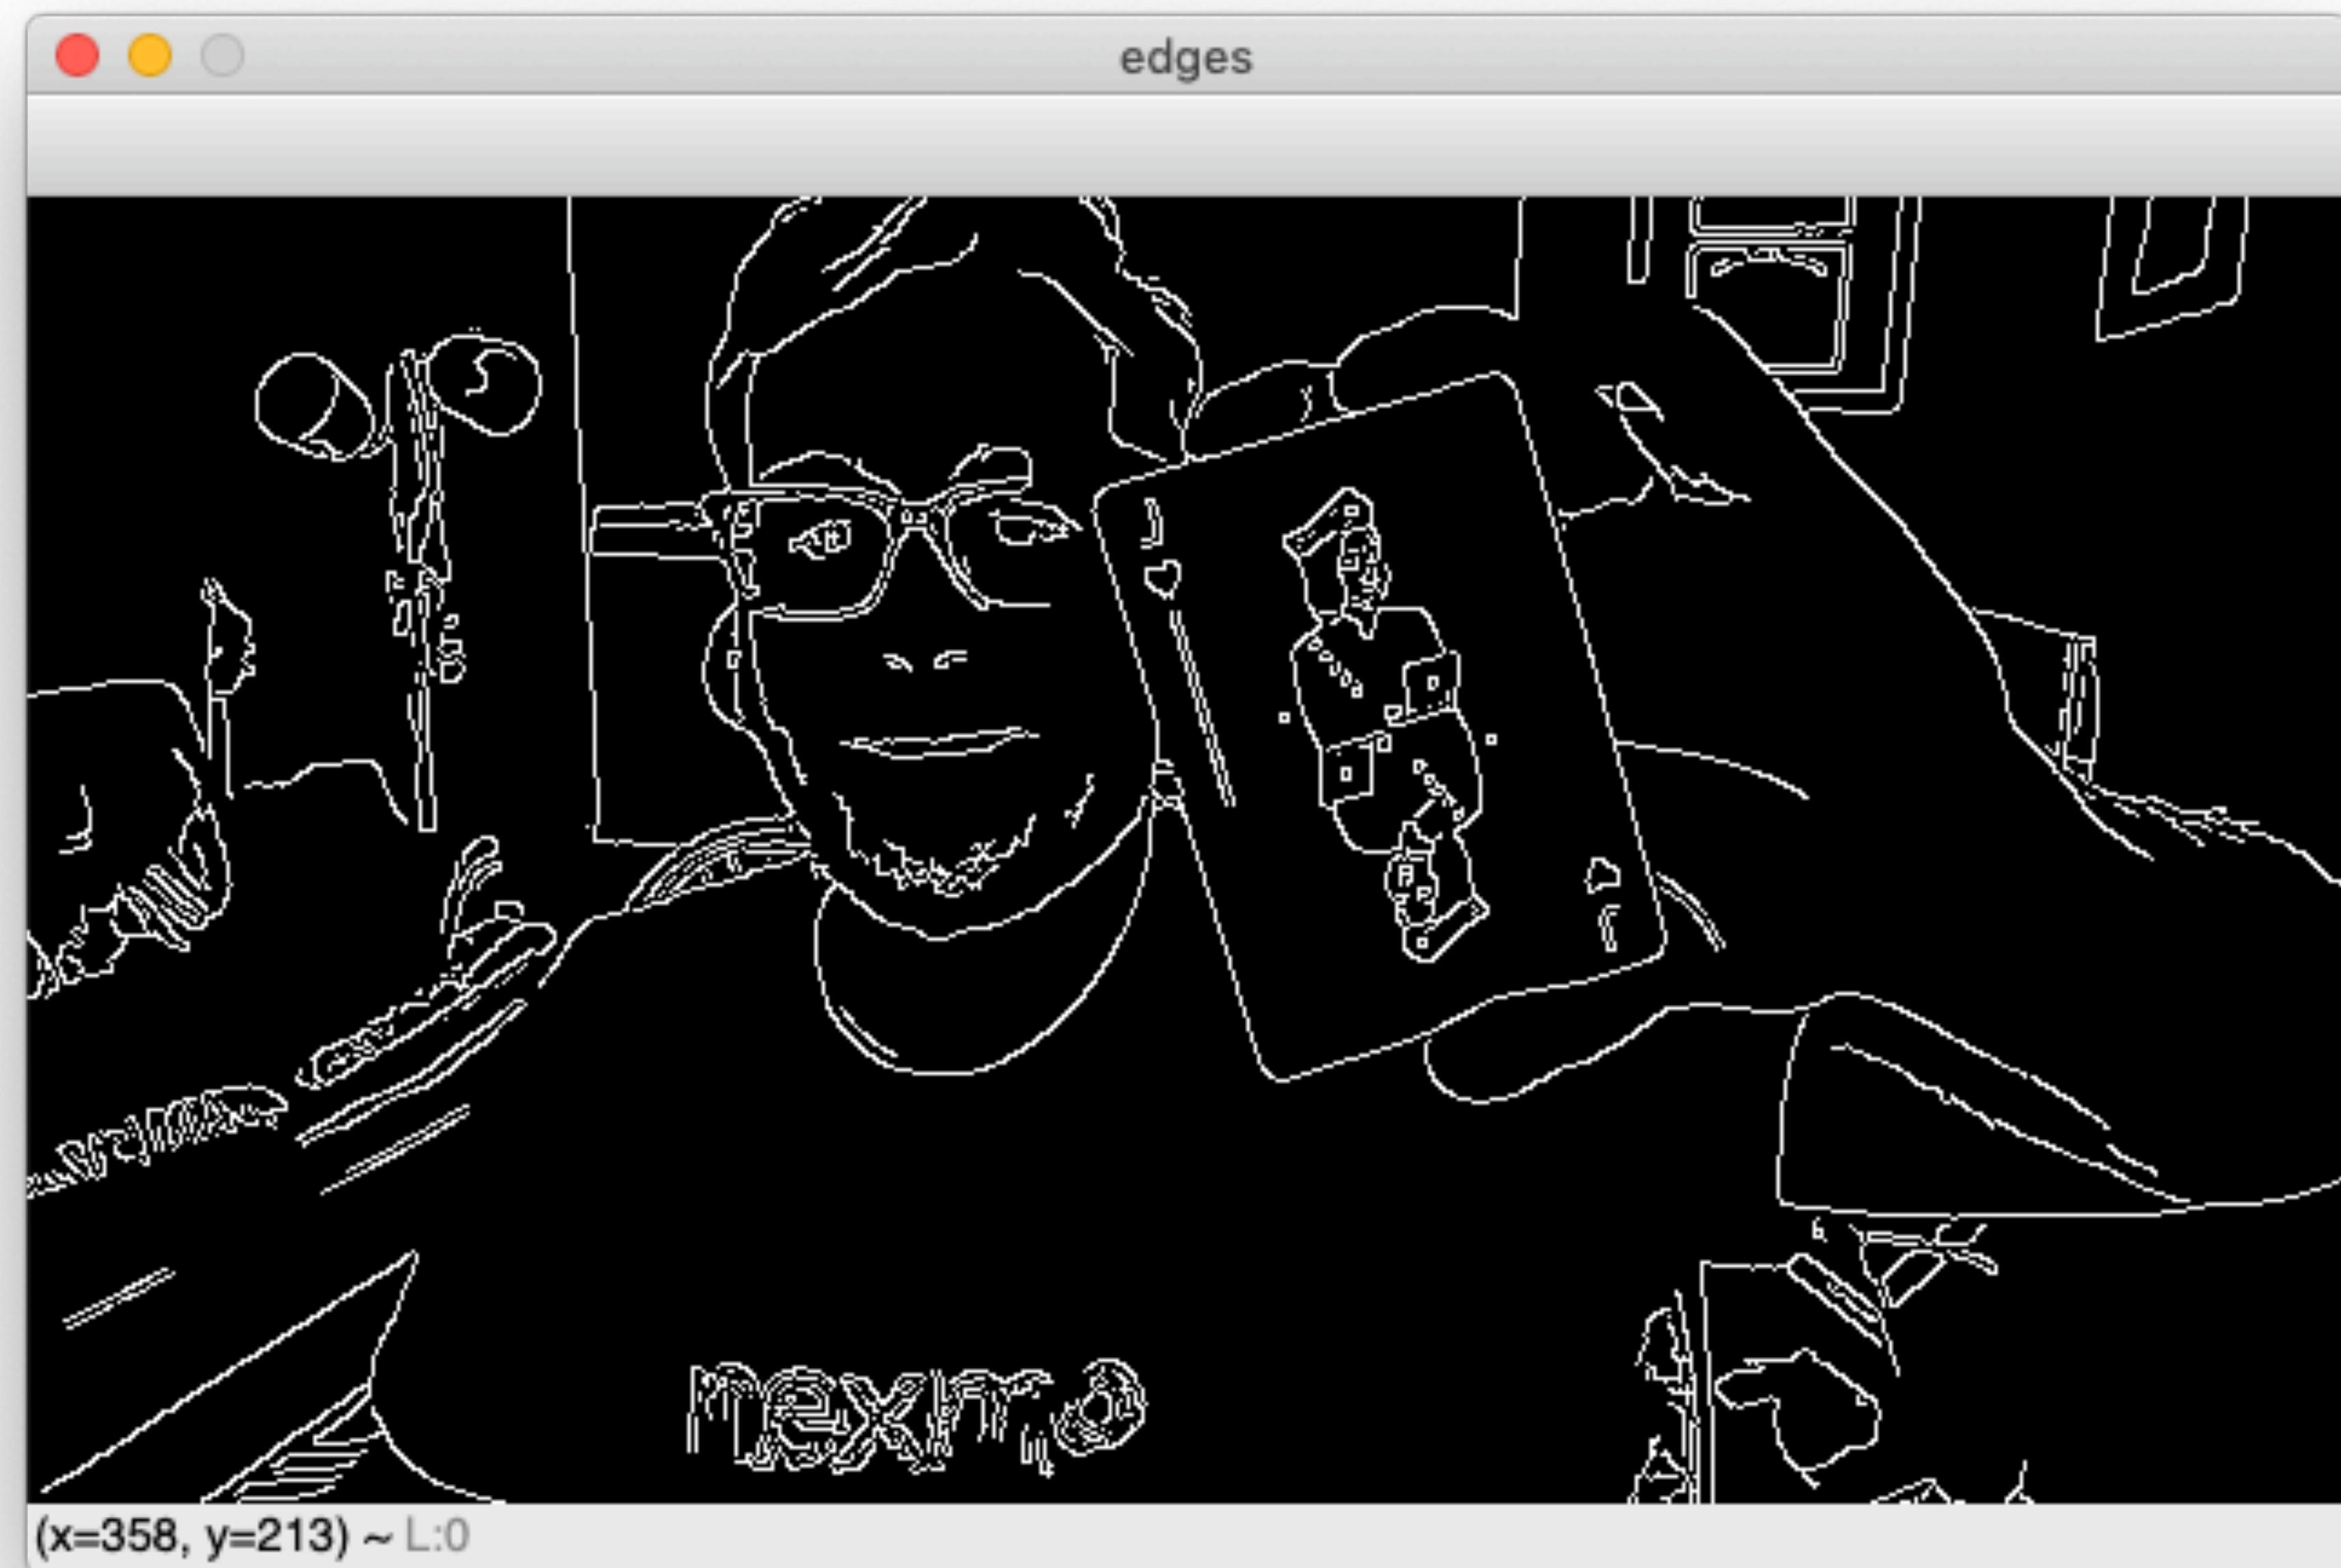
return dilate
```

Maximum Threshold

Minimum Threshold



@cr0wst

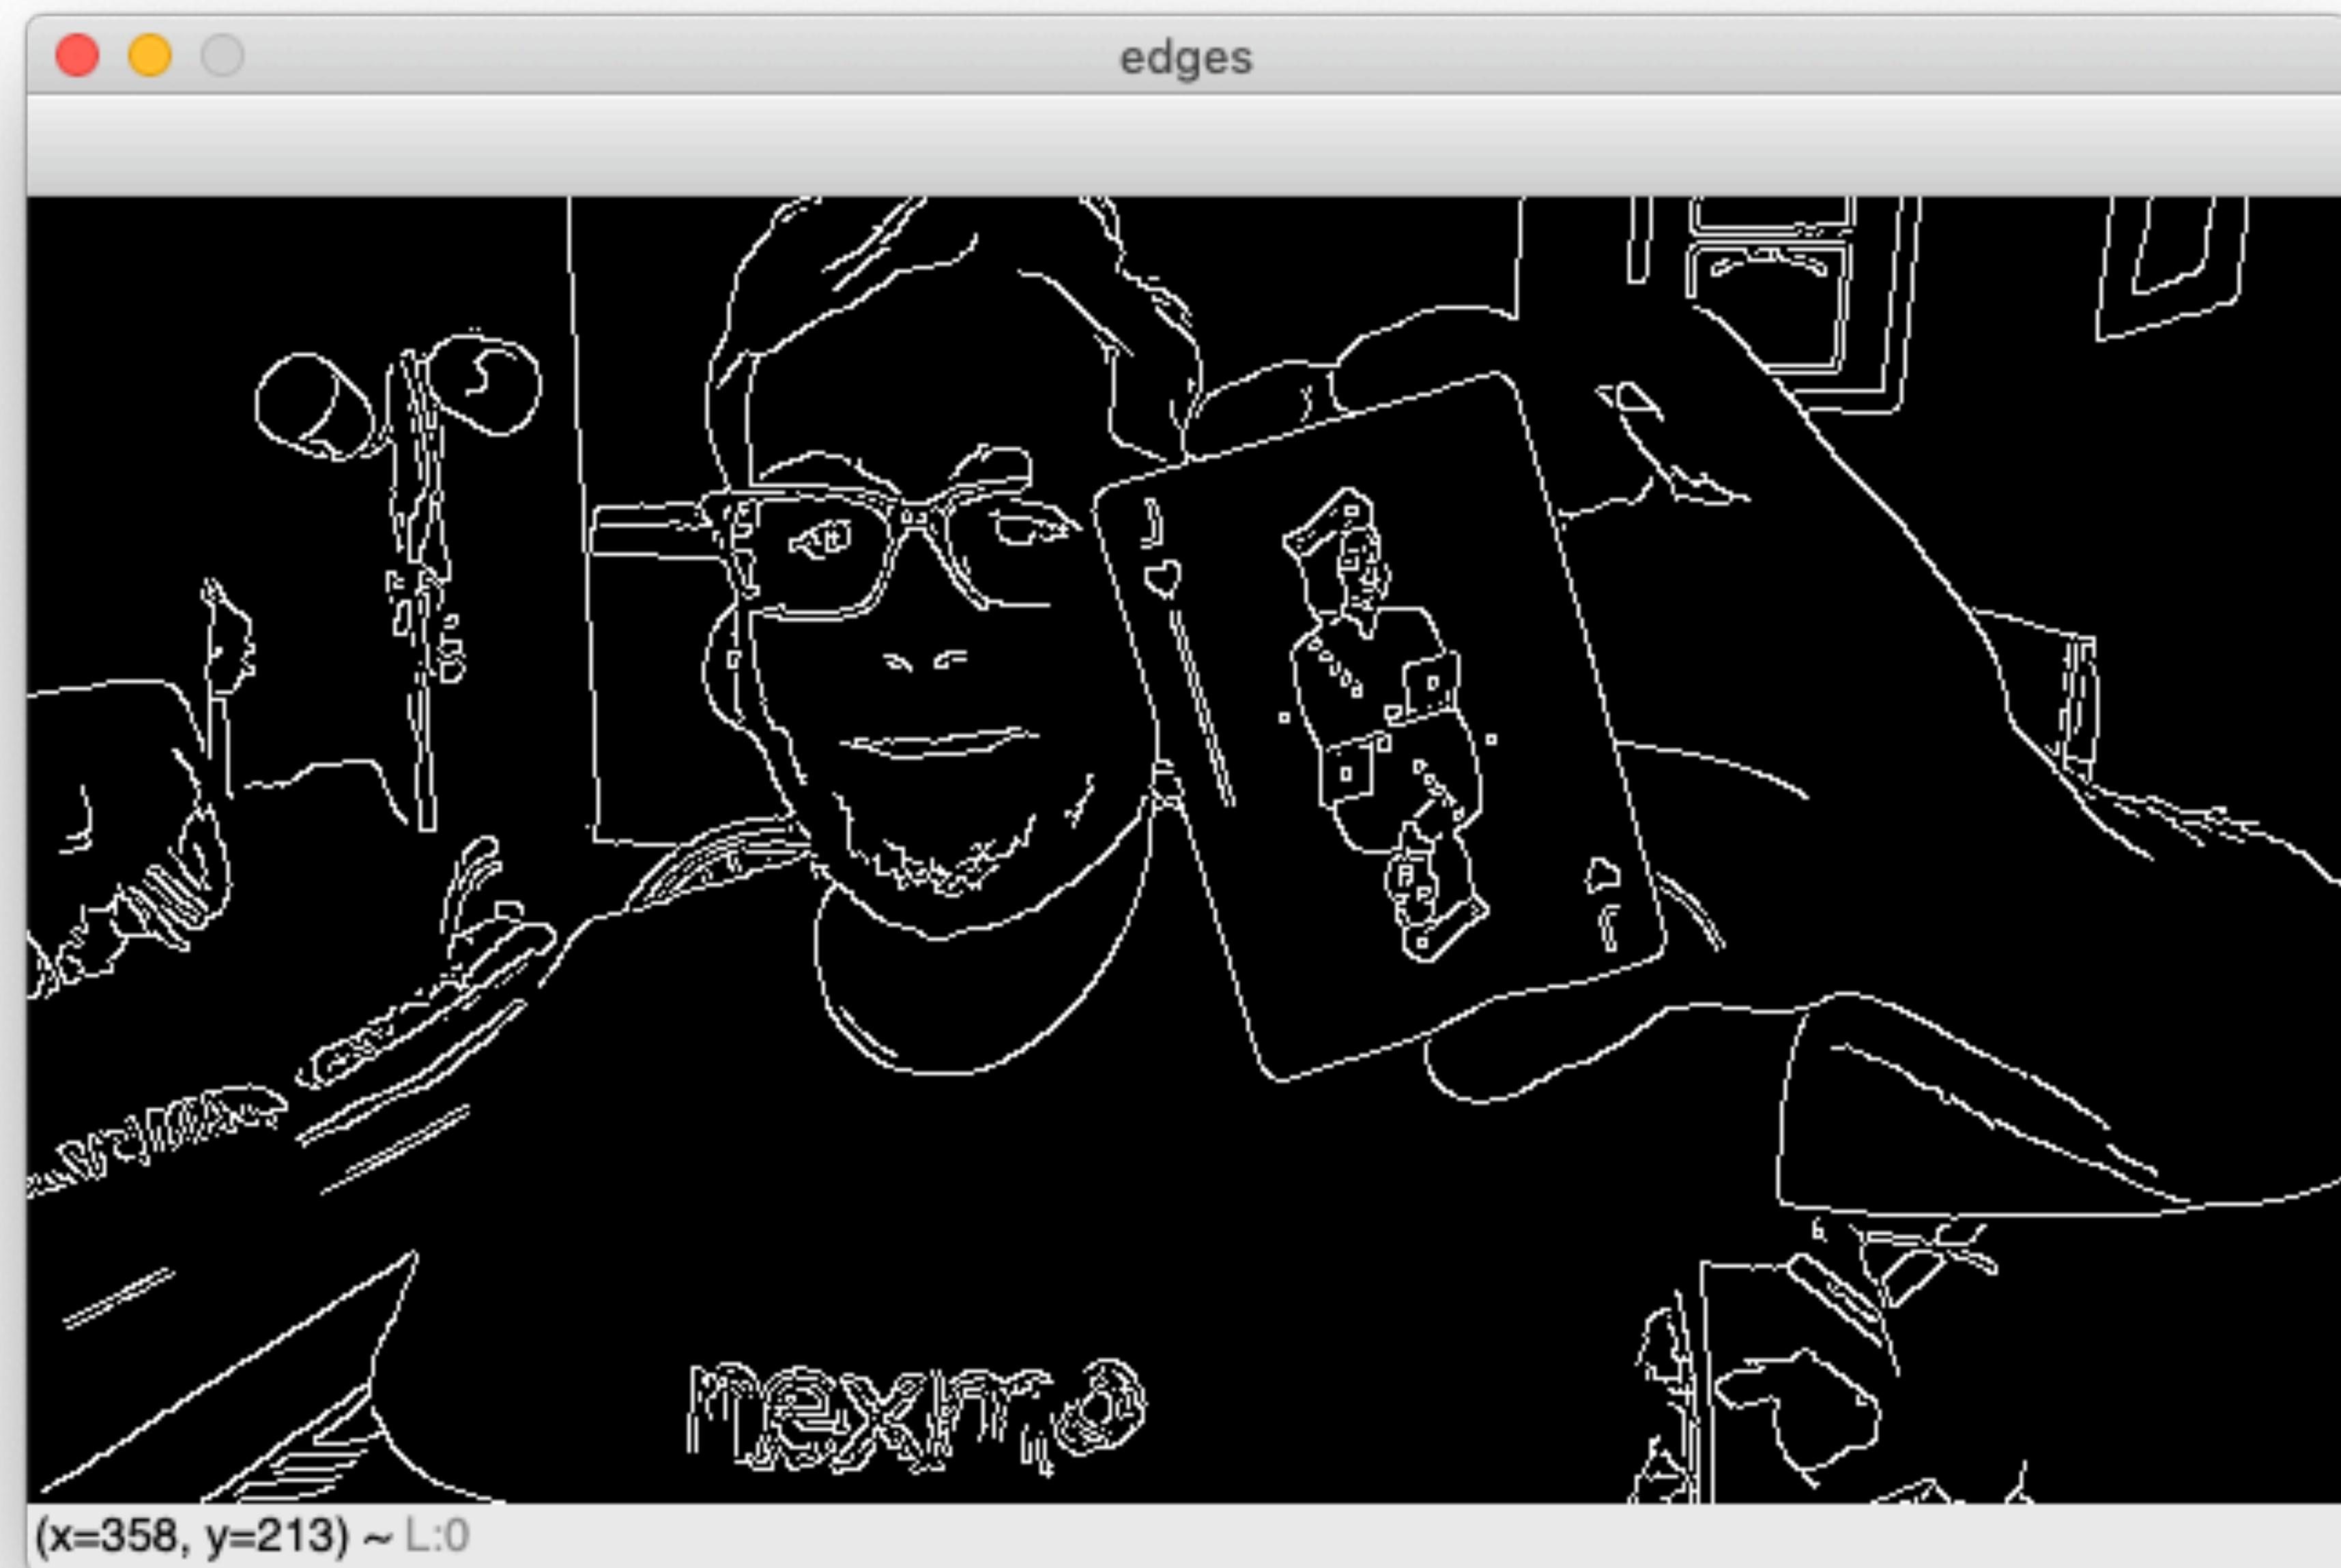


@cr0wst

Why Canny Edge Detection?



@cr0wst



@cr0wst

```
import numpy as np

def normalize(frame):
    frame = imutils.resize(frame, 640)
    gray = cv2.cvtColor(frame, cv2.COLOR_BGR2GRAY)
    blur = cv2.bilateralFilter(gray, 10, 15, 15)

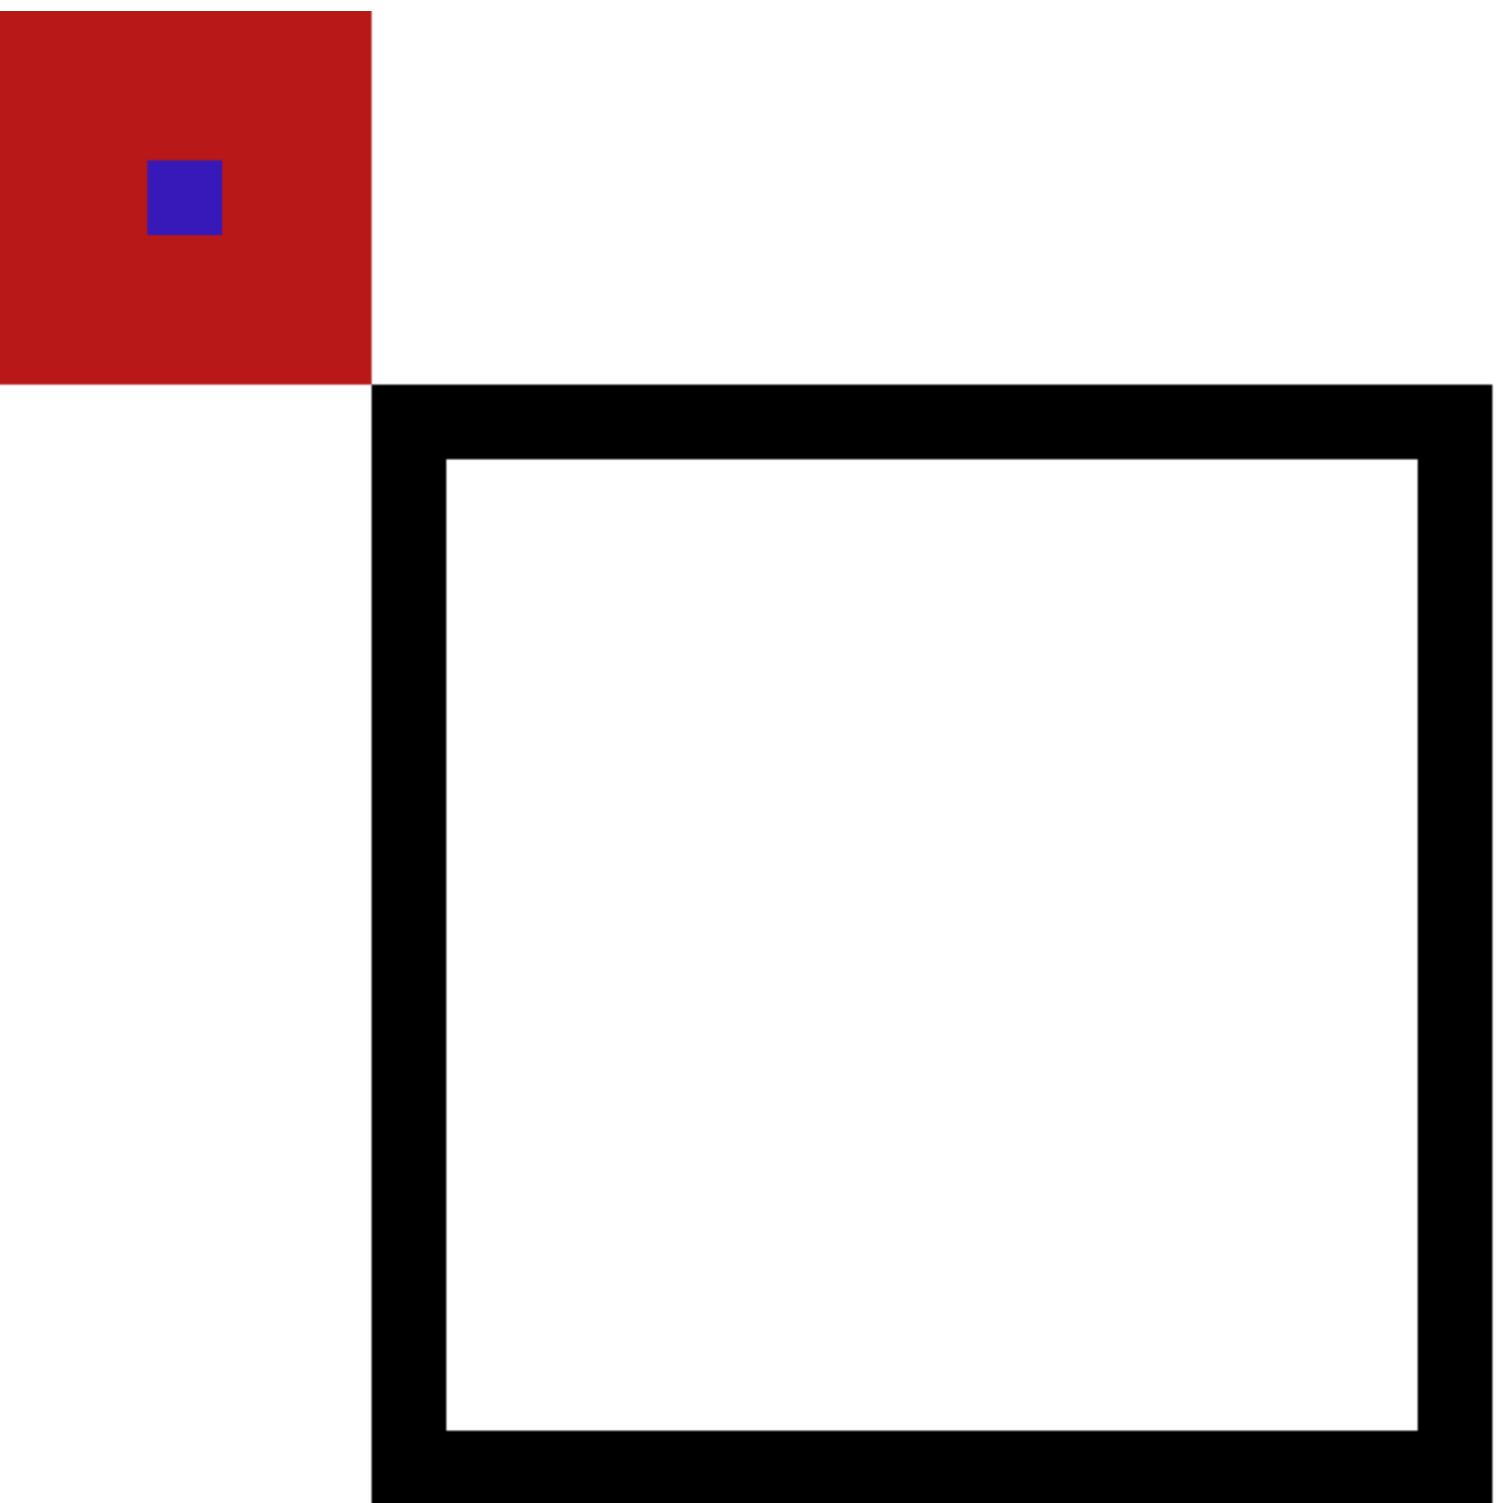
kernel = np.ones((5, 5), np.uint8)
dilate = cv2.dilate(edges, kernel, iterations=1)
edges = cv2.Canny(blur, 50, 150, True)
kernel = np.ones((5, 5), np.uint8)
dilate = cv2.dilate(edges, kernel, iterations=1)

return dilate
```

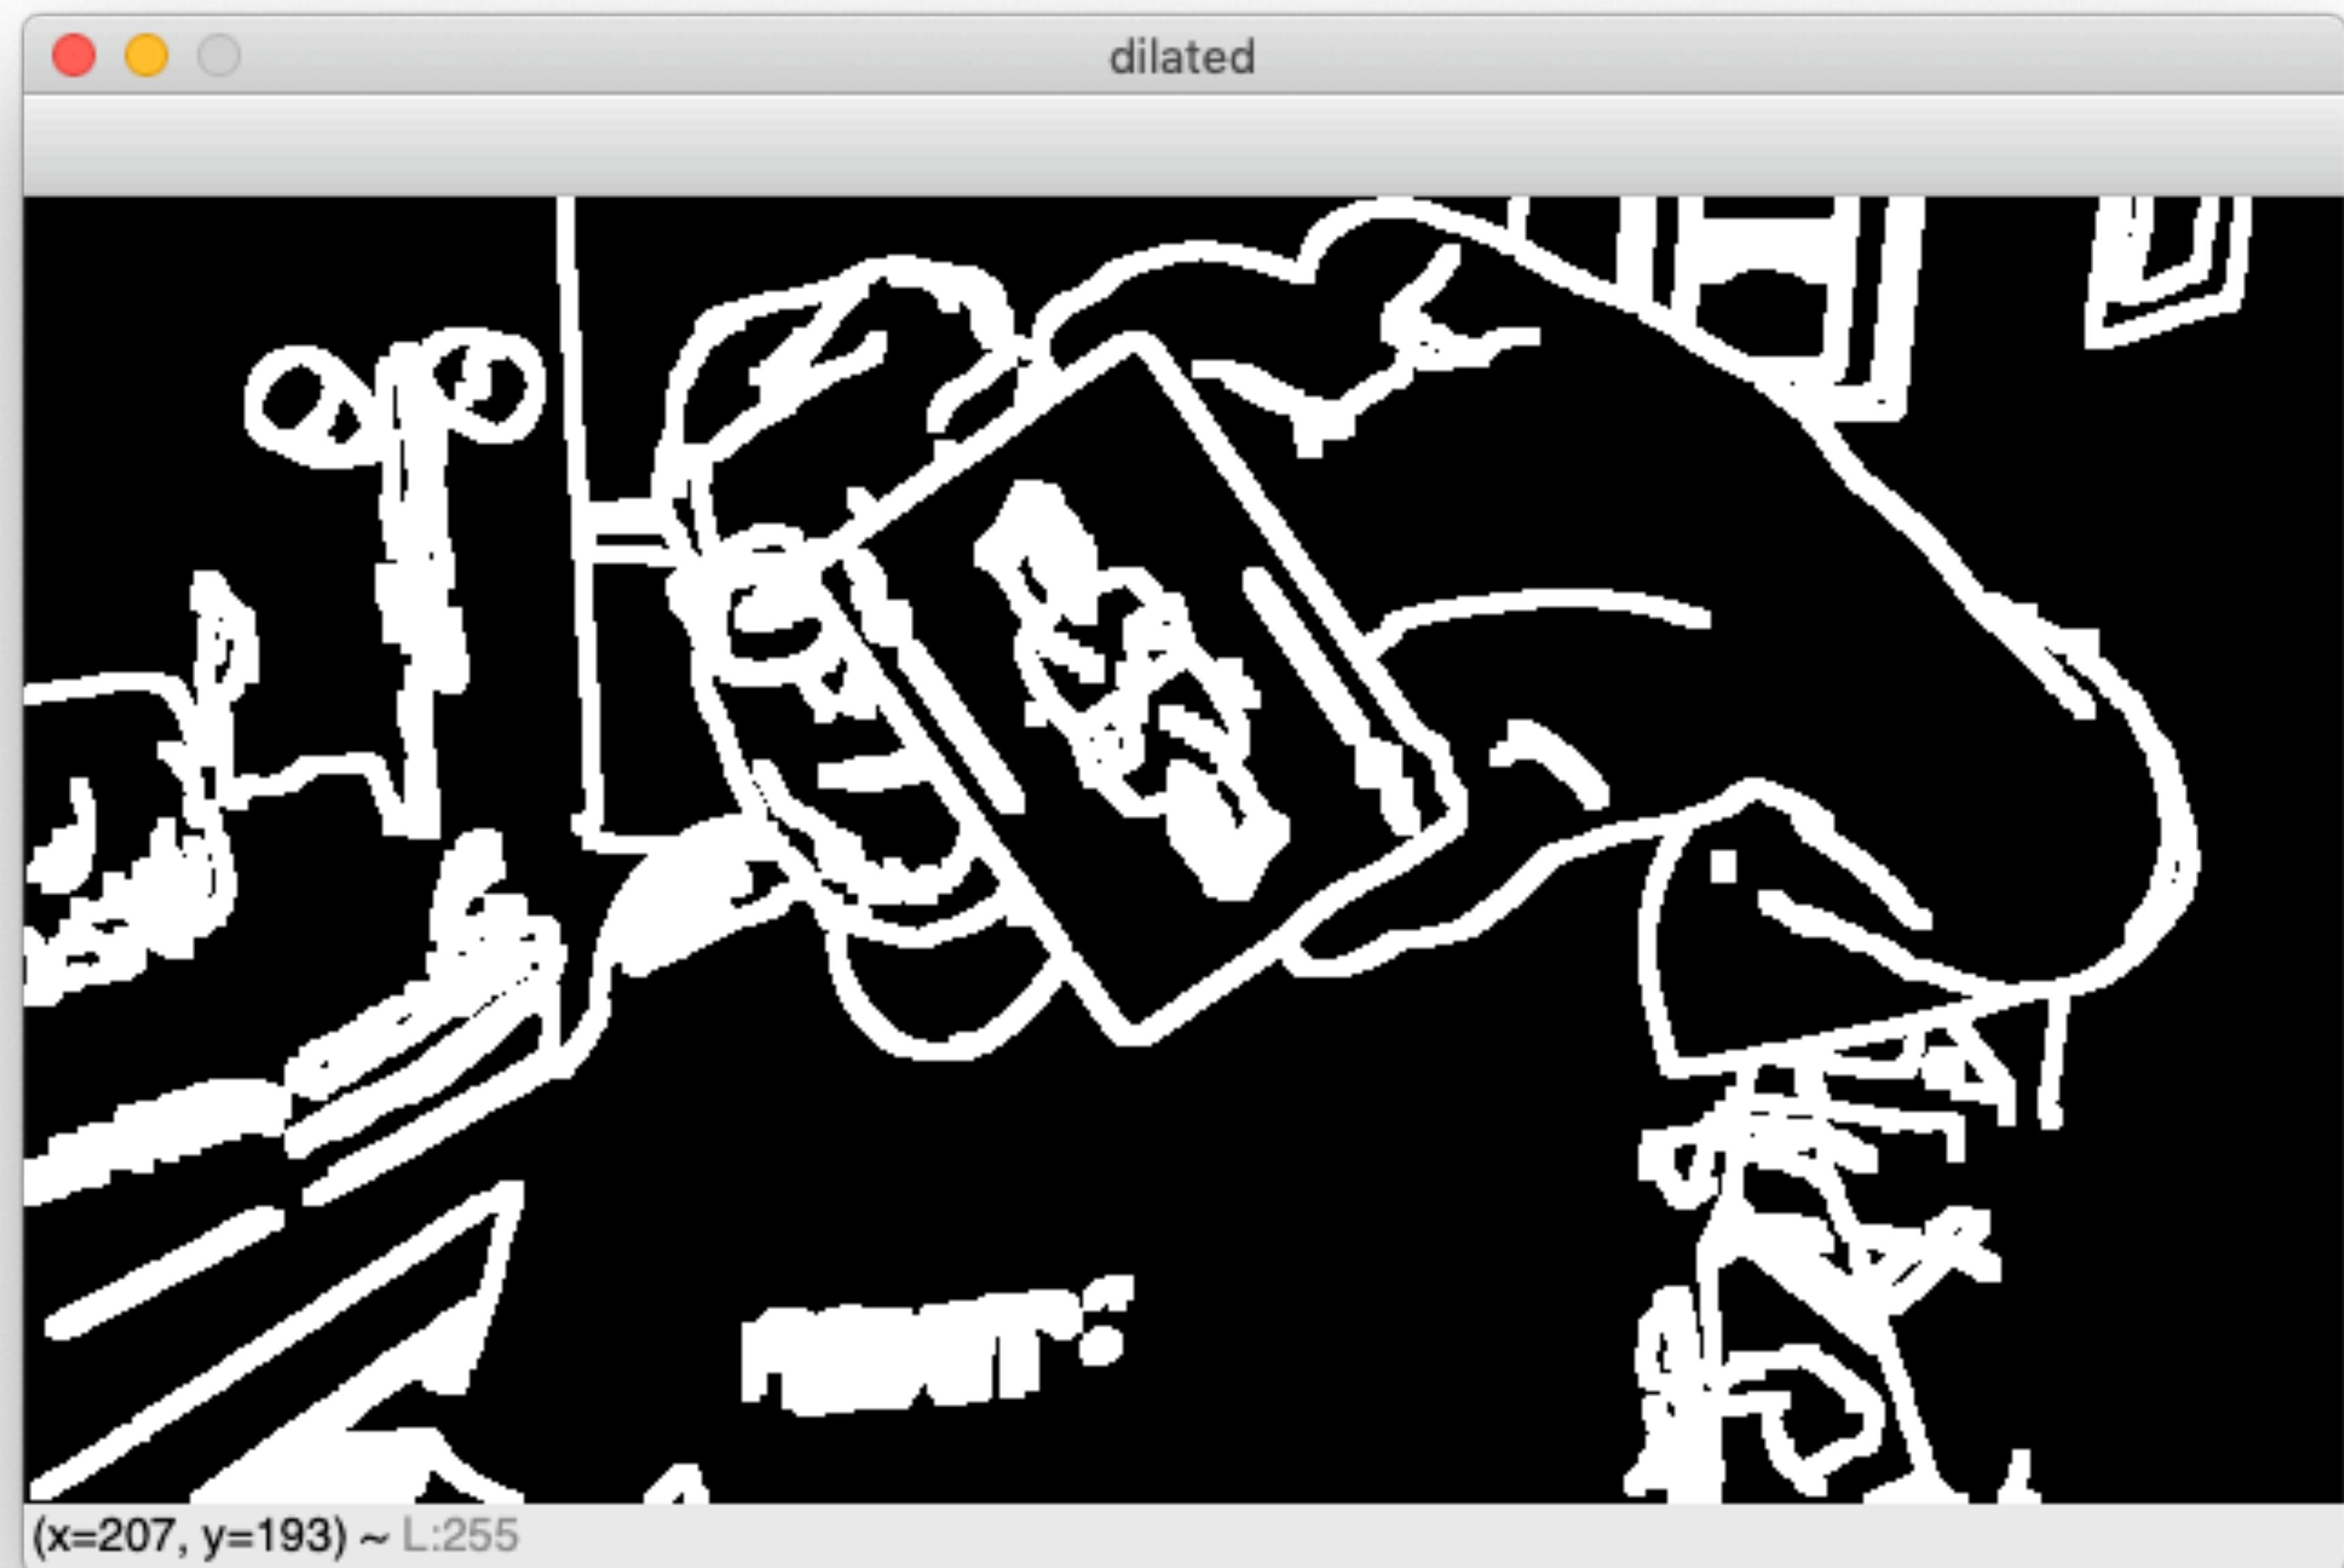
How many times to dilate



@cr0wst



@cr0wst



(x=207, y=193) ~ L:255



@cr0wst

Cards are kind of rectangular



@cr0wst



@cr0wst

```
def get_contours(dilate):
    contours = cv2.findContours(dilate, cv2.RETR_LIST, cv2.CHAIN_APPROX_SIMPLE)
    contours = imutils.grab_contours(contours)
    return sorted(contours, key=cv2.contourArea, reverse=True)
```



@cr0wst

The approximation algorithm. This has to do with how the points are stored.

```
def get_contours(dilate):  
    contours = cv2.findContours(dilate, cv2.RETR_LIST, cv2.CHAIN_APPROX_SIMPLE)  
    contours = imutils.grab_contours(contours)  
    return sorted(contours, key=cv2.contourArea, reverse=True)
```



The retrieval mode. This has to do with edge hierarchies.



@cr0wst

```
def get_contours(dilate):
    contours = cv2.findContours(dilate, cv2.RETR_LIST, cv2.CHAIN_APPROX_SIMPLE)
    contours = imutils.grab_contours(contours)
    return sorted(contours, key=cv2.contourArea, reverse=True)
```



@cr0wst

Remember, rectangles are what we want.



@cr0wst

```
def process(frame):
    image = frame.copy()
    dilate = normalize(frame)
    contours = get_contours(dilate)

    for c in contours:
        rect = cv2.minAreaRect(c)
        box = cv2.boxPoints(rect)
        box = np.int0(box)

        cv2.drawContours(image, [box], -1, (0, 255, 0), 3)
```



@cr0wst

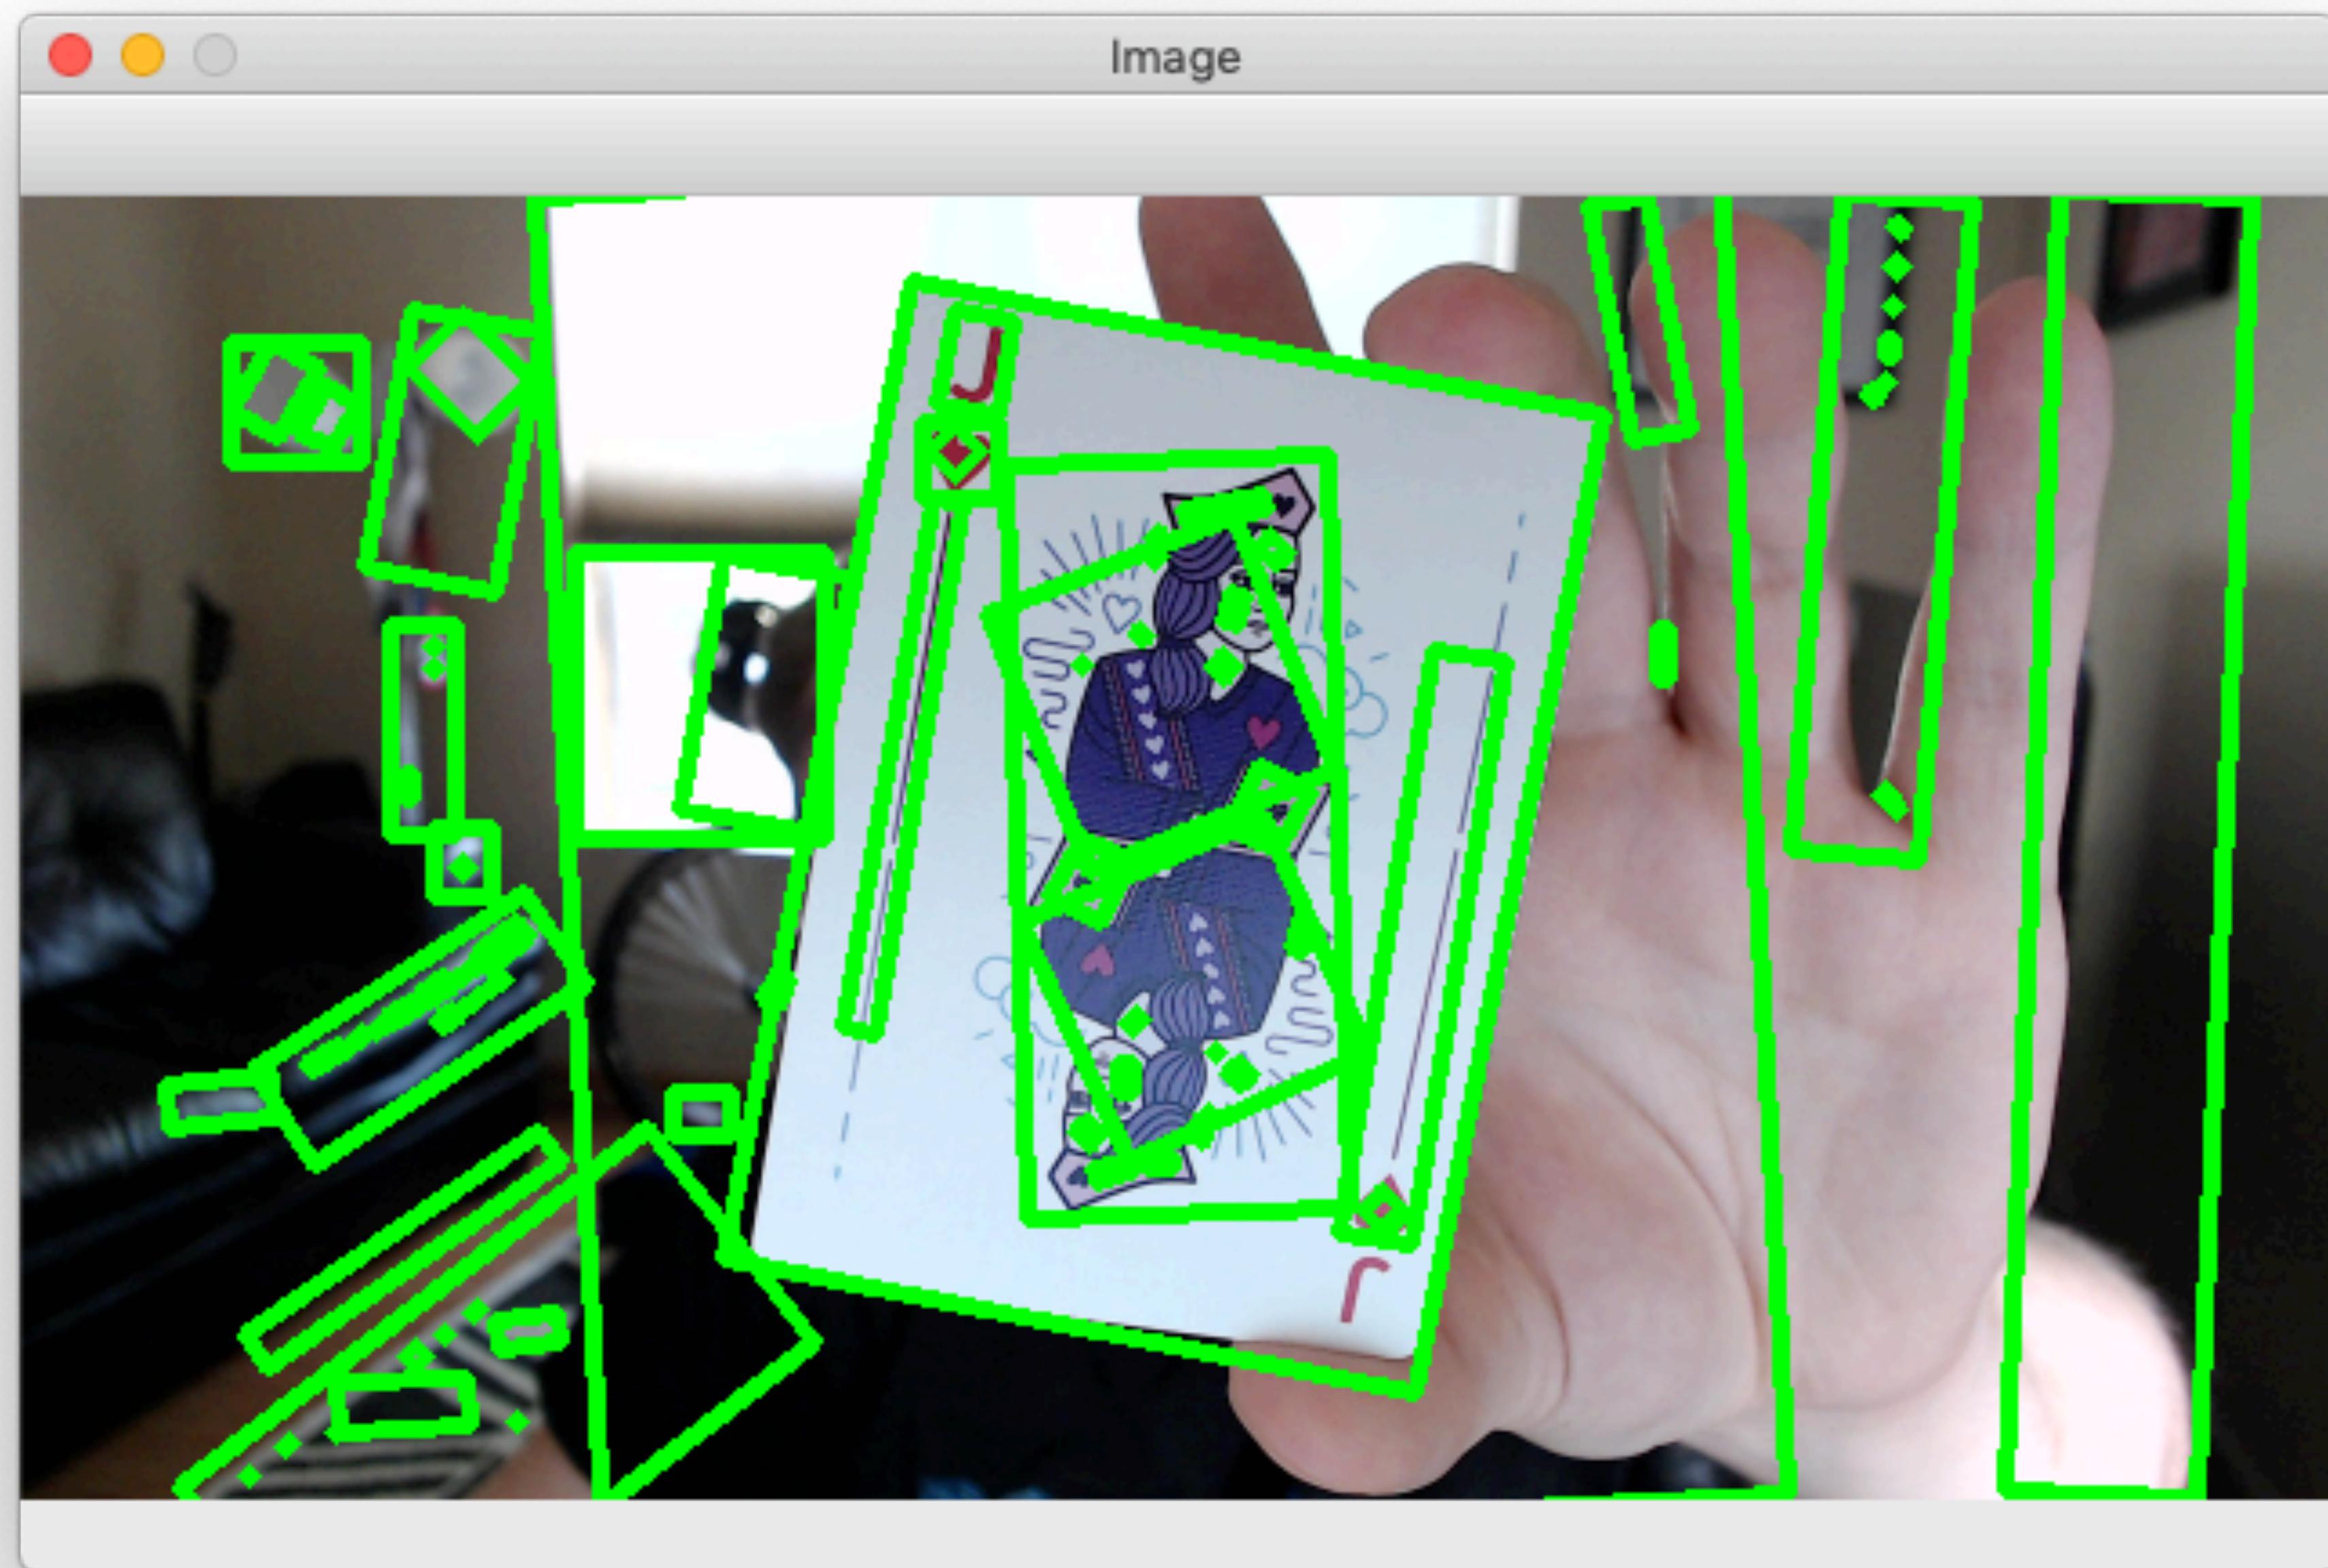
```
def process(frame):
    image = frame.copy()
    dilate = normalize(frame)
    contours = get_contours(dilate)

    for c in contours:
        rect = cv2.minAreaRect(c)
        rect = cv2.minAreaRect(c)
        box = cv2.boxPoints(rect)
        box = np.int0(box)

        cv2.drawContours(image, [box], -1, (0, 255, 0), 3)
```



@cr0wst



@cr0wst

My playing cards have an aspect ratio of around 1.43



@cr0wst

```
((x, y), (w, h), a) = rect = cv2.minAreaRect(c)
box = cv2.boxPoints(rect)
box = np.int0(box)

# sometimes w and h are reversed for turned images
ar = w / float(h) if w > h > 0 else h / float(w)
if 1.40 <= ar <= 1.45:
    cv2.drawContours(image, [box], -1, (0, 255, 0), 3)
```



@cr0wst

```
((x, y), (w, h), a) = rect = cv2.minAreaRect(c)
box = cv2.boxPoints(rect)
box = np.int0(box)

((x, y), (w, h), a) = rect = cv2.minAreaRect(c)
# sometimes w and h are reversed for turned images
ar = w / float(h) if w > h > 0 else h / float(w)
if 1.40 <= ar <= 1.45:
    cv2.drawContours(image, [box], -1, (0, 255, 0), 3)
```

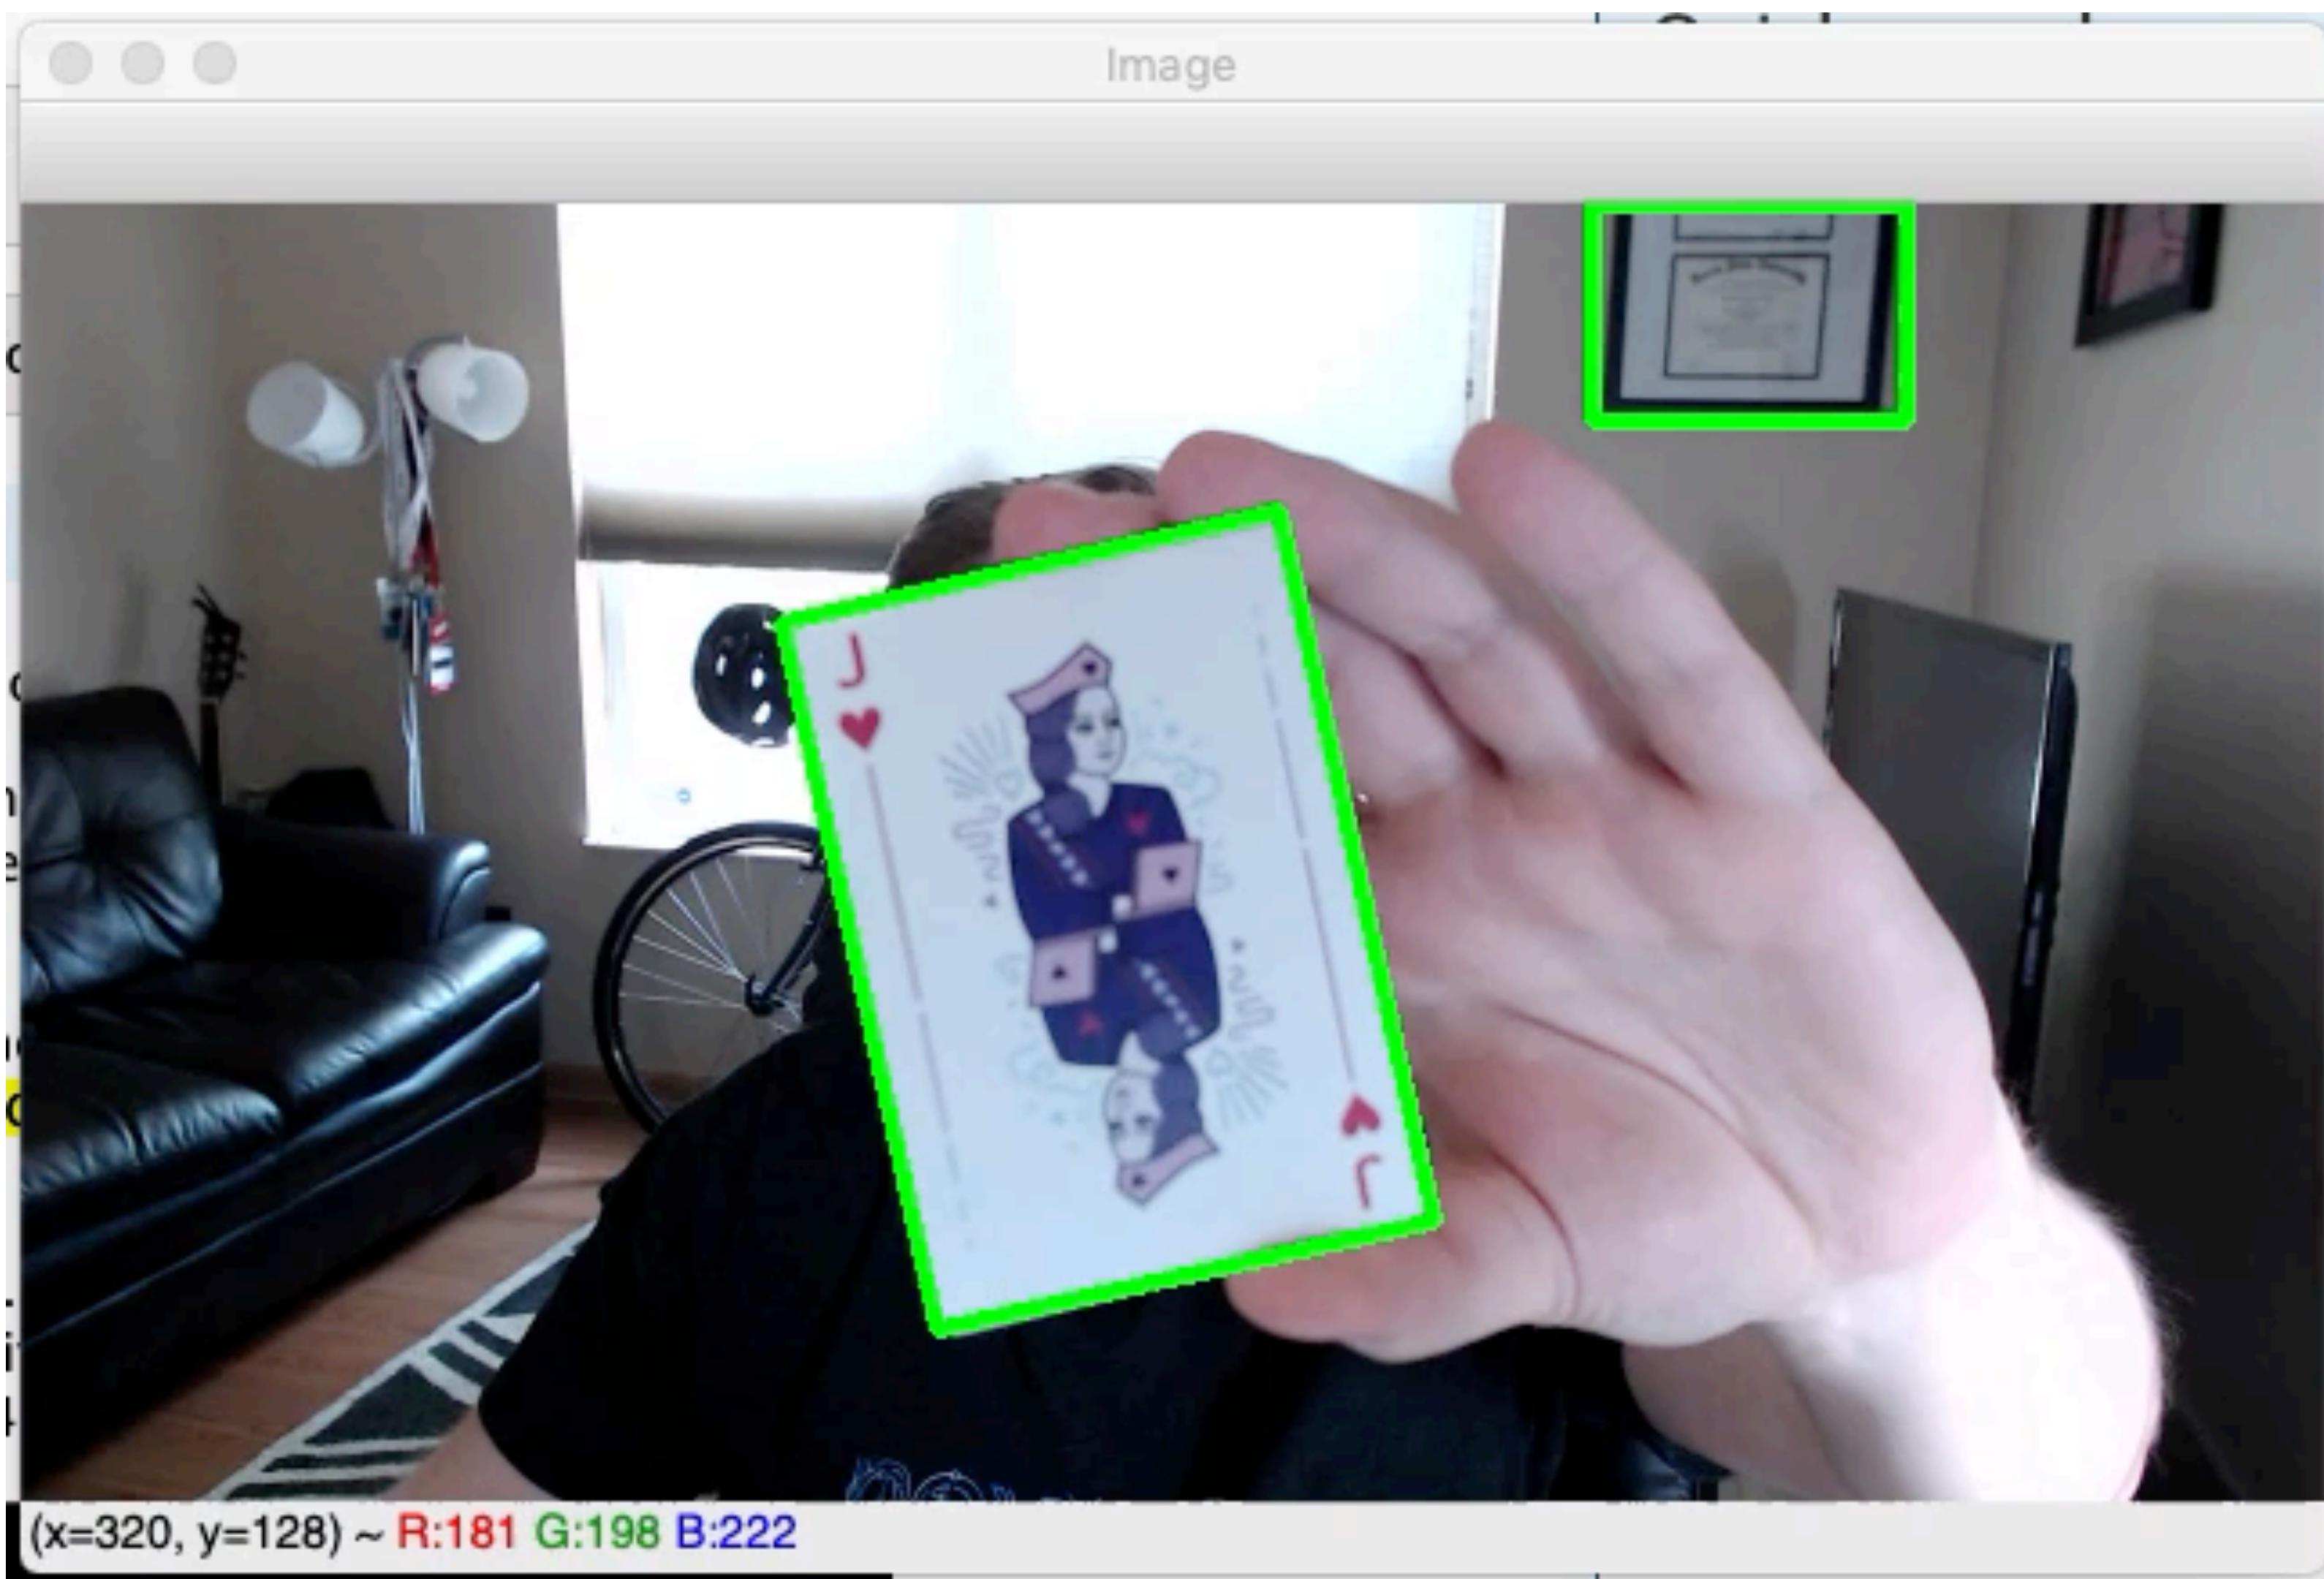


@cr0wst

```
((x, y), (w, h), a) = rect = cv2.minAreaRect(c)
# sometimes w and h are reversed for turned images
box = cv2.boxPoints(rect)
box = np.int0(box)
ar = w / float(h) if w > h > 0 else h / float(w)
if 1.40 <= ar <= 1.45:
    cv2.drawContours(image, [box], -1, (0, 255, 0), 3)
    cv2.drawContours(image, [box], -1, (0, 255, 0), 3)
```



@cr0wst



@cr0wst

```
for c in contours[:2]:  
    ((x, y), (w, h), a) = rect = cv2.minAreaRect(c)  
    box = cv2.boxPoints(rect)  
    box = np.int0(box)  
  
    ar = w / float(h) if w > h > 0 else h / float(w)  
    if 1.40 <= ar <= 1.45:  
        cv2.drawContours(image, [box], -1, (0, 255, 0), 3)  
  
cv2.imshow("Image", image)
```



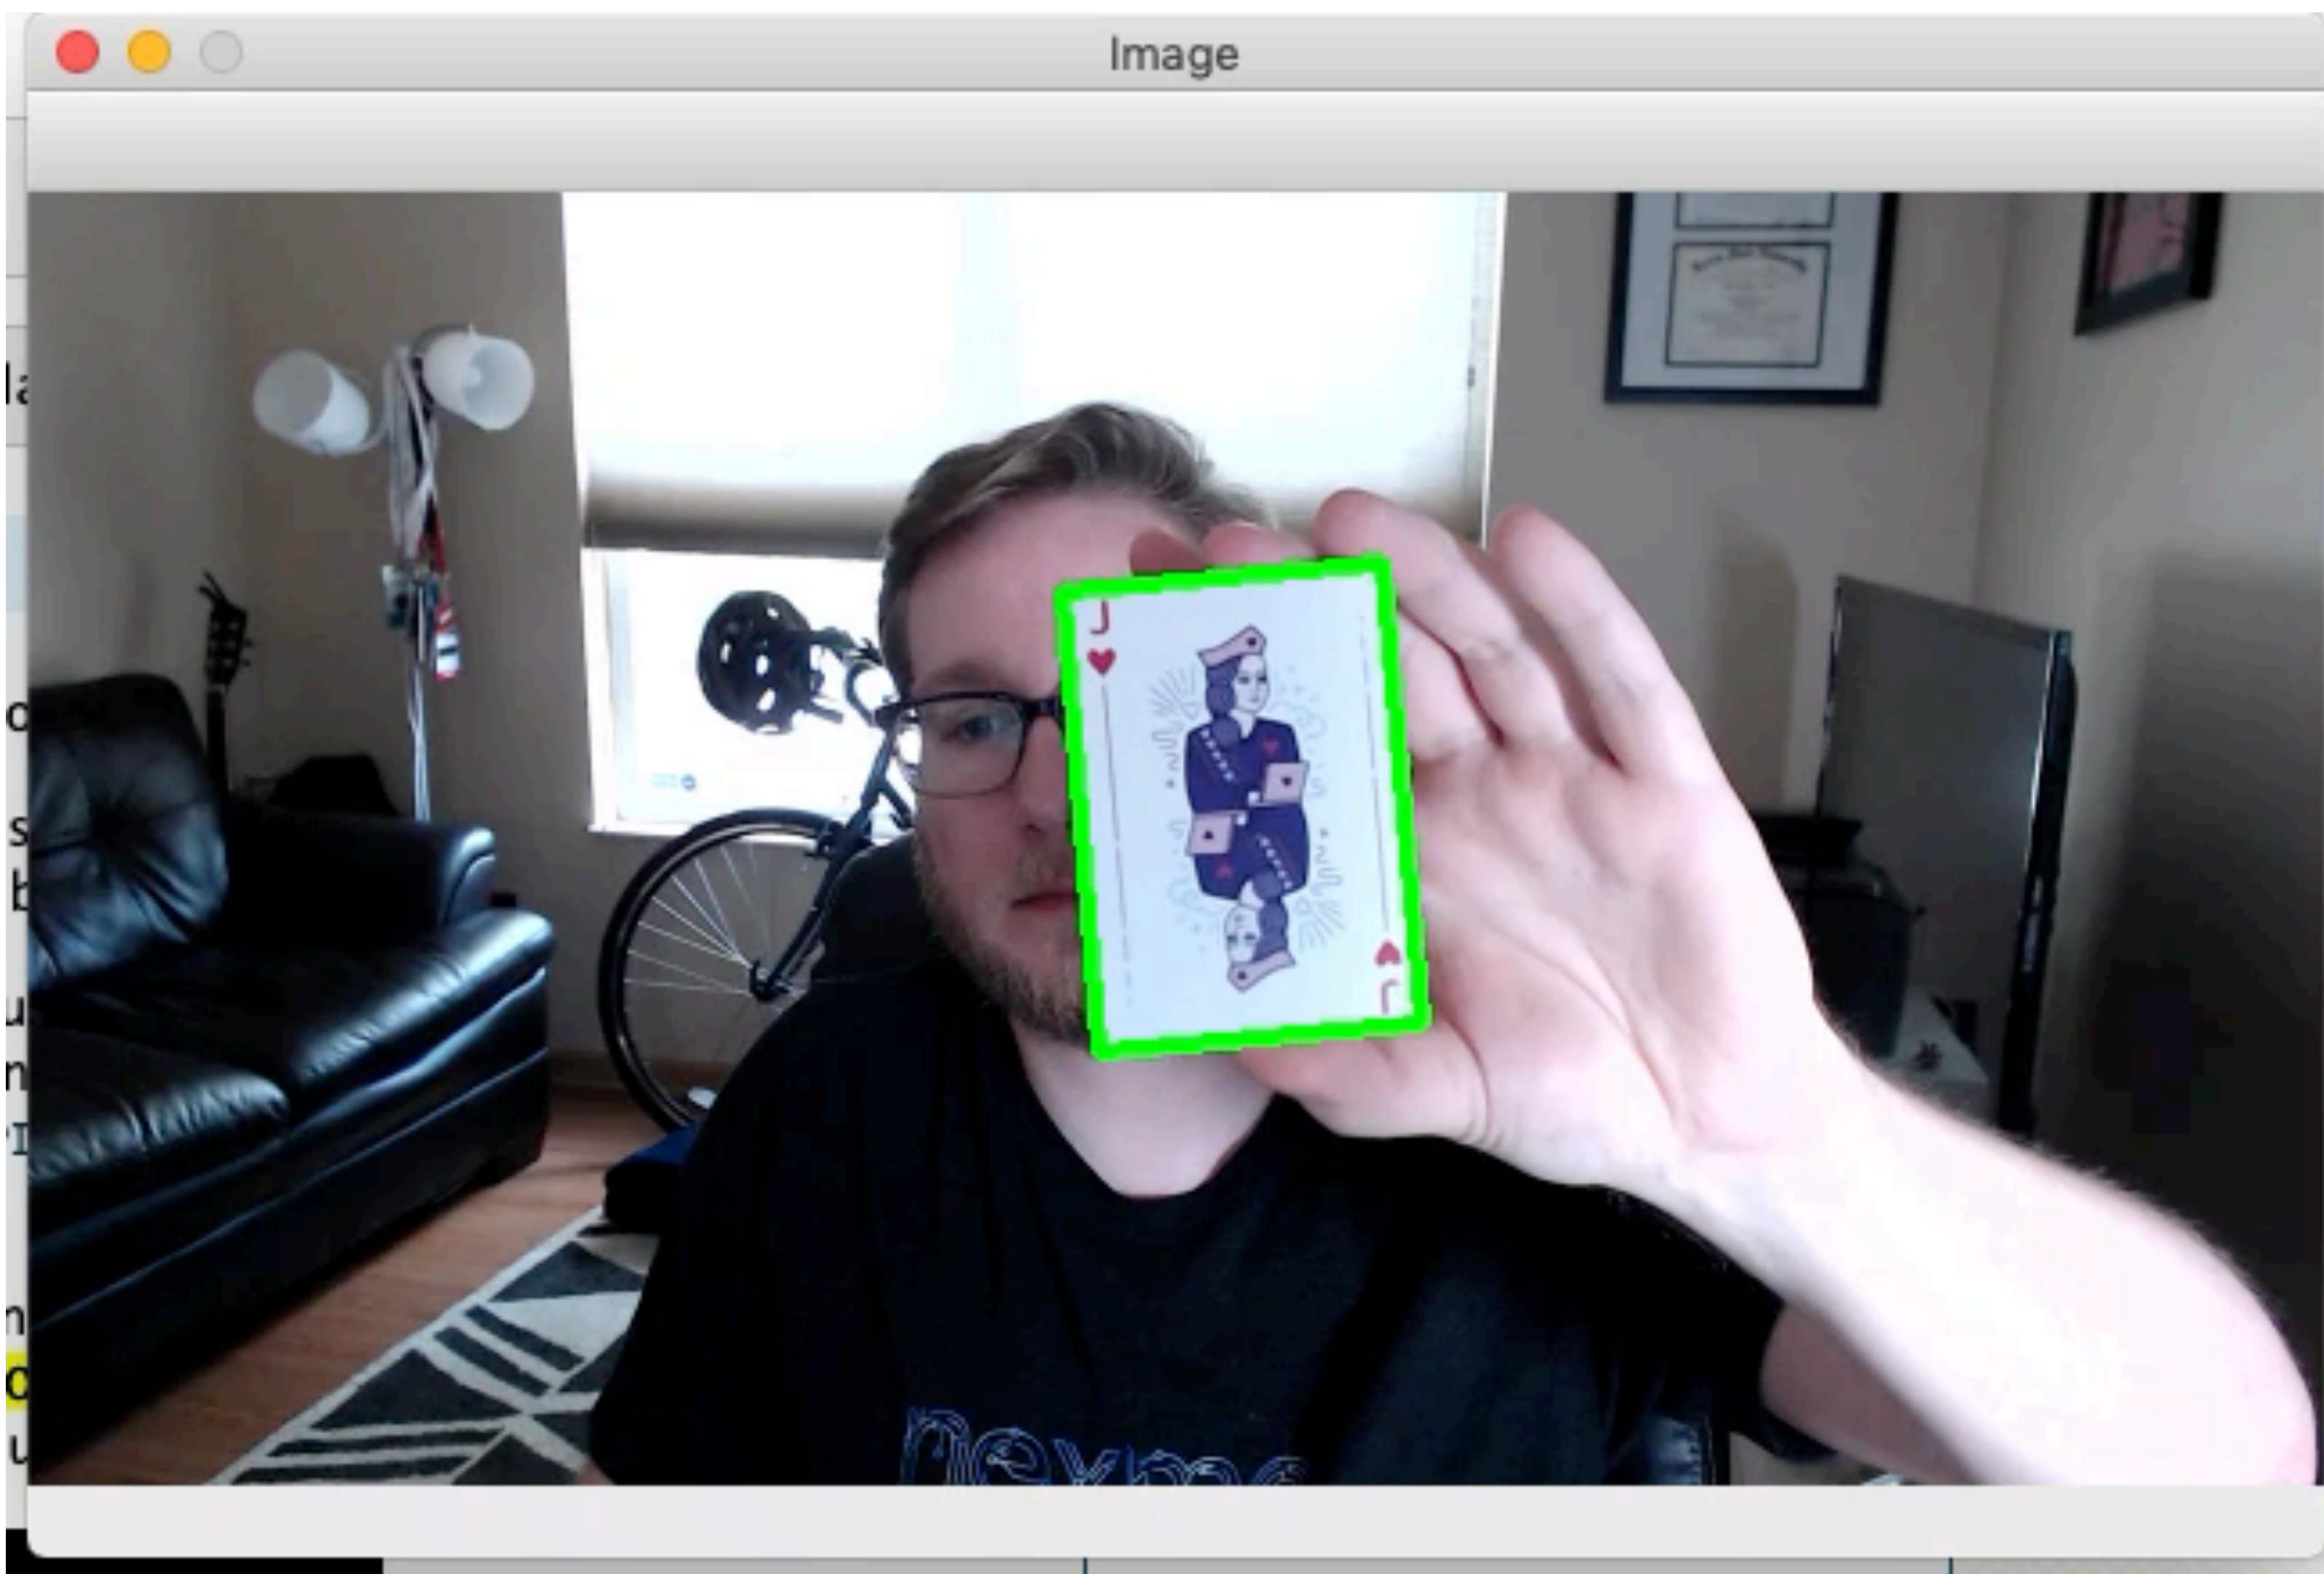
@cr0wst

```
for c in contours[2:]:
    ((x, y), (w, h), a) = rect = cv2.minAreaRect(c)
    box = cv2.boxPoints(rect)
    box = np.int0(box)
    for c in contours[2:]:
        ar = w / float(h) if w > h > 0 else h / float(w)
        if 1.40 <= ar <= 1.45:
            cv2.drawContours(image, [box], -1, (0, 255, 0), 3)

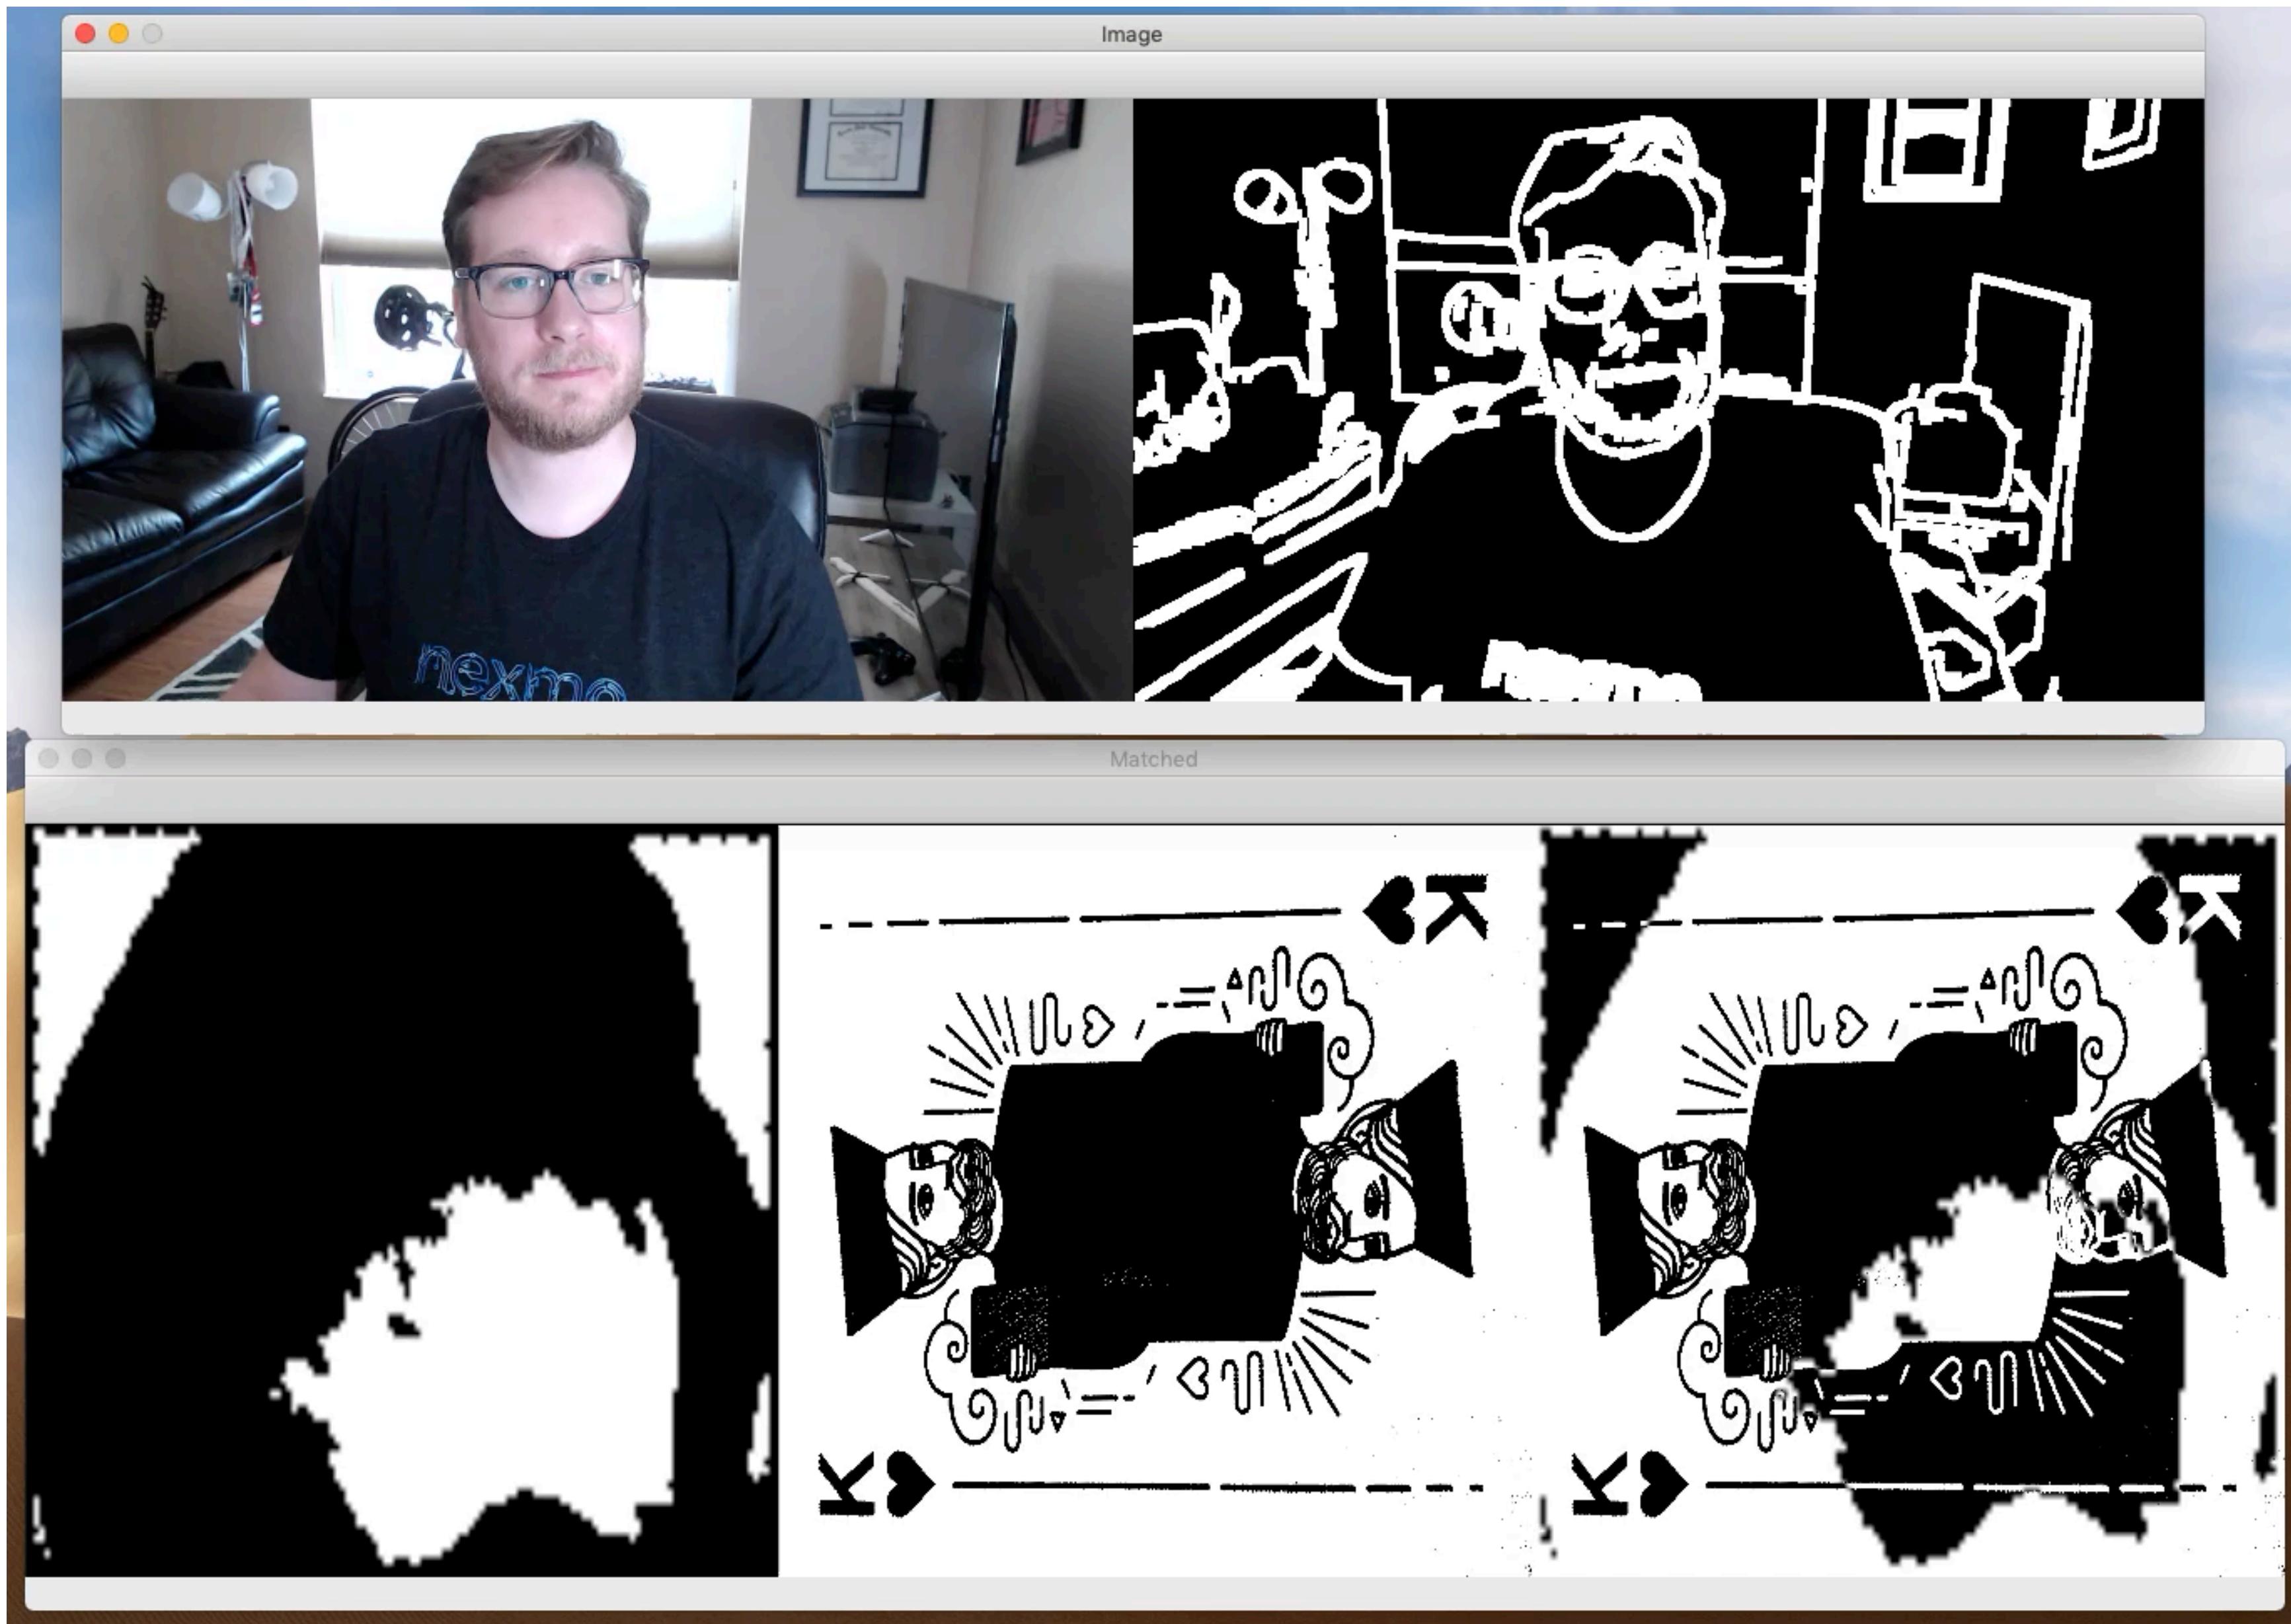
cv2.imshow("Image", image)
```



@cr0wst



@cr0wst



@cr0wst

Simplify when possible.



@cr0wst

Dive in



@cr0wst

I'm Steve!

@cr0wst

That's a zero

steve@smcrow.net



@cr0wst



<https://cr0w.st/cardvision>



@cr0wst